

Ungraded Lab: Training a Sarcasm Detection Model using a Convolution Layer

You will be doing the same steps here as the previous lab but will be using a convolution layer instead. As usual, try tweaking the parameters and observe how it affects the results.

```
In [1]: import json
import matplotlib.pyplot as plt
import tensorflow as tf
```

Load the Dataset

```
In [2]: # The dataset is already downloaded for you. For downloading you can use the code below.
# !wget https://storage.googleapis.com/tensorflow-1-public/course3/sarcasm.json
```

```
In [3]: # Load the JSON file
with open("./sarcasm.json", 'r') as f:
    datastore = json.load(f)

# Initialize the lists
sentences = []
labels = []

# Collect sentences and labels into the lists
for item in datastore:
    sentences.append(item['headline'])
    labels.append(item['is_sarcastic'])
```

Parameters

```
In [4]: # Number of examples to use for training
TRAINING_SIZE = 20000

# Vocabulary size of the tokenizer
VOCAB_SIZE = 10000

# Maximum length of the padded sequences
MAX_LENGTH = 32

# Type of padding
PADDING_TYPE = 'pre'

# Specifies how to truncate the sequences
TRUNC_TYPE = 'post'
```

Split the Dataset

```
In [5]: # Split the sentences
train_sentences = sentences[0:TRAINING_SIZE]
test_sentences = sentences[TRAINING_SIZE:]

# Split the labels
train_labels = labels[0:TRAINING_SIZE]
test_labels = labels[TRAINING_SIZE:]
```

Data preprocessing

```
In [6]: # Instantiate the vectorization layer
vectorize_layer = tf.keras.layers.TextVectorization(max_tokens=VOCAB_SIZE)

# Generate the vocabulary based on the training inputs
vectorize_layer.adapt(train_sentences)
```

```
In [7]: # Preprocess the train and test data
train_dataset = tf.data.Dataset.from_tensor_slices((train_sentences, train_labels))
test_dataset = tf.data.Dataset.from_tensor_slices((test_sentences, test_labels))
```

```
In [8]: def preprocessing_fn(dataset):  
    '''Generates padded sequences from a tf.data.Dataset'''  
  
    # Apply the vectorization layer to the reviews  
    dataset_sequences = dataset.map(lambda review, label: (vectorize_layer(review), label))  
  
    # Put all elements in a single ragged batch  
    dataset_sequences = dataset_sequences.ragged_batch(batch_size=dataset_sequences.cardinality())  
  
    # Output a tensor from the single batch. Extract the sequences and labels.  
    sequences, labels = dataset_sequences.get_single_element()  
  
    # Pad the sequences  
    padded_sequences = tf.keras.utils.pad_sequences(sequences.numpy(), maxlen=MAX_LENGTH, truncating=TRUNC_TYPE, padding=PADDING)  
  
    # Convert back to a tf.data.Dataset  
    padded_sequences = tf.data.Dataset.from_tensor_slices(padded_sequences)  
    labels = tf.data.Dataset.from_tensor_slices(labels)  
  
    # Combine the padded sequences and labels  
    dataset_vectorized = tf.data.Dataset.zip((padded_sequences, labels))  
  
    return dataset_vectorized
```

```
In [9]: # Preprocess the train and test data  
train_dataset_vectorized = train_dataset.apply(preprocessing_fn)  
test_dataset_vectorized = test_dataset.apply(preprocessing_fn)
```

```
In [10]: # View 2 training sequences and its labels  
for example in train_dataset_vectorized.take(2):  
    print(example)  
    print()
```

```
(<tf.Tensor: shape=(32,), dtype=int32, numpy=  
array([ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  
        0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  
       943, 4079, 2366,  47, 366,  94, 2026,  6, 2653, 9470],  
      dtype=int32)>, <tf.Tensor: shape=(), dtype=int32, numpy=0>)
```

```
(<tf.Tensor: shape=(32,), dtype=int32, numpy=  
array([ 0,  0,  0,  0,  0,  0,  0,  0,  0,  0,  
        0,  0,  0,  0,  0,  0,  0,  4, 7185, 3128, 3305,  
       28,  2, 152,  1, 358, 2902,  6, 236,  9, 844],  
      dtype=int32)>, <tf.Tensor: shape=(), dtype=int32, numpy=0>)
```

```
In [11]: SHUFFLE_BUFFER_SIZE = 1000  
PREFETCH_BUFFER_SIZE = tf.data.AUTOTUNE  
BATCH_SIZE = 32  
  
# Optimize and batch the datasets for training  
train_dataset_final = (train_dataset_vectorized  
                        .cache()  
                        .shuffle(SHUFFLE_BUFFER_SIZE)  
                        .prefetch(PREFETCH_BUFFER_SIZE)  
                        .batch(BATCH_SIZE)  
                        )  
  
test_dataset_final = (test_dataset_vectorized  
                      .cache()  
                      .prefetch(PREFETCH_BUFFER_SIZE)  
                      .batch(BATCH_SIZE)  
                      )
```

```
In [12]: def plot_loss_acc(history):
    '''Plots the training and validation loss and accuracy from a history object'''
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs = range(len(acc))

    fig, ax = plt.subplots(1,2, figsize=(12, 6))
    ax[0].plot(epochs, acc, 'bo', label='Training accuracy')
    ax[0].plot(epochs, val_acc, 'b', label='Validation accuracy')
```

```

ax[0].set_title('Training and validation accuracy')
ax[0].set_xlabel('epochs')
ax[0].set_ylabel('accuracy')
ax[0].legend()

ax[1].plot(epochs, loss, 'bo', label='Training Loss')
ax[1].plot(epochs, val_loss, 'b', label='Validation Loss')
ax[1].set_title('Training and validation loss')
ax[1].set_xlabel('epochs')
ax[1].set_ylabel('loss')
ax[1].legend()

plt.show()

```

Build and Compile the Model

```

In [13]: # Parameters
EMBEDDING_DIM = 16
FILTERS = 128
KERNEL_SIZE = 5
DENSE_DIM = 6

# Model Definition with Conv1D
model_conv = tf.keras.Sequential([
    tf.keras.Input(shape=(MAX_LENGTH,)),
    tf.keras.layers.Embedding(input_dim=VOCAB_SIZE, output_dim=EMBEDDING_DIM),
    tf.keras.layers.Conv1D(FILTERS, KERNEL_SIZE, activation='relu'),
    tf.keras.layers.GlobalMaxPooling1D(),
    tf.keras.layers.Dense(DENSE_DIM, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

# Set the training parameters
model_conv.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# Print the model summary
model_conv.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 32, 16)	160,000
conv1d (Conv1D)	(None, 28, 128)	10,368
global_max_pooling1d (GlobalMaxPooling1D)	(None, 128)	0
dense (Dense)	(None, 6)	774
dense_1 (Dense)	(None, 1)	7

Total params: 171,149 (668.55 KB)

Trainable params: 171,149 (668.55 KB)

Non-trainable params: 0 (0.00 B)

Train the Model

```

In [14]: NUM_EPOCHS = 10

# Train the model
history_conv = model_conv.fit(train_dataset_final, epochs=NUM_EPOCHS, validation_data=test_dataset_final)

```

Epoch 1/10

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1745365816.757086 13117 service.cc:145] XLA service 0x76f80400d000 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1745365816.757125 13117 service.cc:153] StreamExecutor device (0): NVIDIA A10G, Compute Capability 8.6

157/625 0s 972us/step - accuracy: 0.5650 - loss: 0.6715

I0000 00:00:1745365818.306757 13117 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

```

625/625 — 4s 2ms/step - accuracy: 0.6932 - loss: 0.5398 - val_accuracy: 0.8419 - val_loss: 0.3462
Epoch 2/10
625/625 — 1s 1ms/step - accuracy: 0.8961 - loss: 0.2537 - val_accuracy: 0.8471 - val_loss: 0.3802
Epoch 3/10
625/625 — 1s 1ms/step - accuracy: 0.9502 - loss: 0.1476 - val_accuracy: 0.8410 - val_loss: 0.4816
Epoch 4/10
625/625 — 1s 1ms/step - accuracy: 0.9770 - loss: 0.0752 - val_accuracy: 0.8377 - val_loss: 0.5955
Epoch 5/10
625/625 — 1s 1ms/step - accuracy: 0.9916 - loss: 0.0331 - val_accuracy: 0.8402 - val_loss: 0.6927
Epoch 6/10
625/625 — 1s 1ms/step - accuracy: 0.9971 - loss: 0.0129 - val_accuracy: 0.8322 - val_loss: 0.7782
Epoch 7/10
625/625 — 1s 1ms/step - accuracy: 0.9991 - loss: 0.0051 - val_accuracy: 0.8307 - val_loss: 0.8715
Epoch 8/10
625/625 — 1s 1ms/step - accuracy: 0.9997 - loss: 0.0021 - val_accuracy: 0.8278 - val_loss: 0.9698
Epoch 9/10
625/625 — 1s 1ms/step - accuracy: 0.9999 - loss: 9.1340e-04 - val_accuracy: 0.8283 - val_loss: 1.0911
Epoch 10/10
625/625 — 1s 1ms/step - accuracy: 0.9999 - loss: 6.4805e-04 - val_accuracy: 0.8258 - val_loss: 1.1334

```

In [15]: `plot_loss_acc(history_conv)`

