

# Ungraded Lab: Using a multi-layer LSTM for forecasting

In this lab, you will use the same RNN architecture in the first lab but will instead stack [LSTM](#) layers instead of `SimpleRNN`.

## Imports

```
In [1]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
```

## Utilities

```
In [2]: def plot_series(time, series, format="-", start=0, end=None):
    """
    Visualizes time series data

    Args:
        time (array of int) - contains the time steps
        series (array of int) - contains the measurements for each time step
        format - line style when plotting the graph
        start - first time step to plot
        end - last time step to plot
    """

    # Setup dimensions of the graph figure
    plt.figure(figsize=(10, 6))

    if type(series) is tuple:

        for series_num in series:
            # Plot the time series data
            plt.plot(time[start:end], series_num[start:end], format)

    else:
        # Plot the time series data
        plt.plot(time[start:end], series[start:end], format)

    # Label the x-axis
    plt.xlabel("Time")

    # Label the y-axis
    plt.ylabel("Value")

    # Overlay a grid on the graph
    plt.grid(True)

    # Draw the graph on screen
    plt.show()

def trend(time, slope=0):
    """
    Generates synthetic data that follows a straight line given a slope value.

    Args:
        time (array of int) - contains the time steps
        slope (float) - determines the direction and steepness of the line

    Returns:
        series (array of float) - measurements that follow a straight line
    """

    # Compute the linear series given the slope
    series = slope * time

    return series

def seasonal_pattern(season_time):
    """
    Just an arbitrary pattern, you can change it if you wish

    Args:
```

```

    season_time (array of float) - contains the measurements per time step

Returns:
    data_pattern (array of float) - contains revised measurement values according
                                    to the defined pattern
"""

# Generate the values using an arbitrary pattern
data_pattern = np.where(season_time < 0.4,
                        np.cos(season_time * 2 * np.pi),
                        1 / np.exp(3 * season_time))

return data_pattern

def seasonality(time, period, amplitude=1, phase=0):
    """
    Repeats the same pattern at each period

    Args:
        time (array of int) - contains the time steps
        period (int) - number of time steps before the pattern repeats
        amplitude (int) - peak measured value in a period
        phase (int) - number of time steps to shift the measured values

    Returns:
        data_pattern (array of float) - seasonal data scaled by the defined amplitude
    """

    # Define the measured values per period
    season_time = ((time + phase) % period) / period

    # Generates the seasonal data scaled by the defined amplitude
    data_pattern = amplitude * seasonal_pattern(season_time)

    return data_pattern

def noise(time, noise_level=1, seed=None):
    """Generates a normally distributed noisy signal

    Args:
        time (array of int) - contains the time steps
        noise_level (float) - scaling factor for the generated signal
        seed (int) - number generator seed for repeatability

    Returns:
        noise (array of float) - the noisy signal
    """

    # Initialize the random number generator
    rnd = np.random.RandomState(seed)

    # Generate a random number for each time step and scale by the noise level
    noise = rnd.randn(len(time)) * noise_level

    return noise

```

## Generate the Synthetic Data

```


In [3]: # Parameters
time = np.arange(4 * 365 + 1, dtype="float32")
baseline = 10
amplitude = 40
slope = 0.05
noise_level = 5

# Create the series
series = baseline + trend(time, slope) + seasonality(time, period=365, amplitude=amplitude)

# Update with noise
series += noise(time, noise_level, seed=42)

# Plot the results
plot_series(time, series)

```



## Split the Dataset

```
In [4]: # Define the split time
split_time = 1000

# Get the train set
time_train = time[:split_time]
x_train = series[:split_time]

# Get the validation set
time_valid = time[split_time:]
x_valid = series[split_time:]
```

## Prepare Features and Labels

```
In [5]: # Parameters
window_size = 20
batch_size = 32
shuffle_buffer_size = 1000
```

```
In [6]: def windowed_dataset(series, window_size, batch_size, shuffle_buffer):
    """Generates dataset windows

    Args:
        series (array of float) - contains the values of the time series
        window_size (int) - the number of time steps to include in the feature
        batch_size (int) - the batch size
        shuffle_buffer(int) - buffer size to use for the shuffle method

    Returns:
        dataset (TF Dataset) - TF Dataset containing time windows
    """

    # Add an axis for the feature dimension of RNN layers
    series = tf.expand_dims(series, axis=-1)

    # Generate a TF Dataset from the series values
    dataset = tf.data.Dataset.from_tensor_slices(series)

    # Window the data but only take those with the specified size
    dataset = dataset.window(window_size + 1, shift=1, drop_remainder=True)

    # Flatten the windows by putting its elements in a single batch
```

```

dataset = dataset.flat_map(lambda window: window.batch(window_size + 1))

# Create tuples with features and labels
dataset = dataset.map(lambda window: (window[:-1], window[-1]))

# Shuffle the windows
dataset = dataset.shuffle(shuffle_buffer)

# Create batches of windows
dataset = dataset.batch(batch_size)

# Optimize the dataset for training
dataset = dataset.cache().prefetch(1)

return dataset

```

```

In [7]: # Generate the dataset windows
dataset = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

```

## Build the Model

As mentioned, you will swap `SimplerNN` for `LSTM` in this lab. It is also set as `bidirectional` below but feel free to revise later and see what results you get. LSTMs are much more complex in their internal architecture than simpleRNNs. It implements a cell state that allows it to remember sequences better than simple implementations. This added complexity results in a bigger set of parameters to train and you'll see that when you print the model summary below.

```

In [8]: # Build the Model
model_tune = tf.keras.models.Sequential([
    tf.keras.Input(shape=(window_size, 1)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 100.0)
])

# Print the model summary
model_tune.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
bidirectional (Bidirectional)	(None, 20, 64)	8,704
bidirectional_1 (Bidirectional)	(None, 64)	24,832
dense (Dense)	(None, 1)	65
lambda (Lambda)	(None, 1)	0

Total params: 33,601 (131.25 KB)  
 Trainable params: 33,601 (131.25 KB)  
 Non-trainable params: 0 (0.00 B)

## Tune the Learning Rate

As usual, you will pick a learning rate by running the tuning code below.

```

In [9]: # Set the Learning rate scheduler
lr_schedule = tf.keras.callbacks.LearningRateScheduler(
    lambda epoch: 1e-8 * 10**(epoch / 20))

































# Initialize the optimizer
optimizer = tf.keras.optimizers.SGD(momentum=0.9)

# Set the training parameters
model_tune.compile(loss=tf.keras.losses.Huber(), optimizer=optimizer)

# Train the model
history = model_tune.fit(dataset, epochs=100, callbacks=[lr_schedule])

```

Epoch 1/100  
31/31 ————— 3s 6ms/step - loss: 25.9677 - learning\_rate: 1.0000e-08  
Epoch 2/100  
31/31 ————— 0s 5ms/step - loss: 25.8166 - learning\_rate: 1.1220e-08  
Epoch 3/100  
31/31 ————— 0s 6ms/step - loss: 25.6226 - learning\_rate: 1.2589e-08  
Epoch 4/100  
31/31 ————— 0s 5ms/step - loss: 25.4159 - learning\_rate: 1.4125e-08  
Epoch 5/100  
31/31 ————— 0s 5ms/step - loss: 25.1914 - learning\_rate: 1.5849e-08  
Epoch 6/100  
31/31 ————— 0s 5ms/step - loss: 24.9346 - learning\_rate: 1.7783e-08  
Epoch 7/100  
31/31 ————— 0s 5ms/step - loss: 24.6494 - learning\_rate: 1.9953e-08  
Epoch 8/100  
31/31 ————— 0s 5ms/step - loss: 24.3334 - learning\_rate: 2.2387e-08  
Epoch 9/100  
31/31 ————— 0s 5ms/step - loss: 23.9843 - learning\_rate: 2.5119e-08  
Epoch 10/100  
31/31 ————— 0s 5ms/step - loss: 23.6082 - learning\_rate: 2.8184e-08  
Epoch 11/100  
31/31 ————— 0s 5ms/step - loss: 23.2080 - learning\_rate: 3.1623e-08  
Epoch 12/100  
31/31 ————— 0s 5ms/step - loss: 22.7789 - learning\_rate: 3.5481e-08  
Epoch 13/100  
31/31 ————— 0s 5ms/step - loss: 22.3378 - learning\_rate: 3.9811e-08  
Epoch 14/100  
31/31 ————— 0s 5ms/step - loss: 21.8734 - learning\_rate: 4.4668e-08  
Epoch 15/100  
31/31 ————— 0s 5ms/step - loss: 21.4225 - learning\_rate: 5.0119e-08  
Epoch 16/100  
31/31 ————— 0s 5ms/step - loss: 20.9776 - learning\_rate: 5.6234e-08  
Epoch 17/100  
31/31 ————— 0s 5ms/step - loss: 20.5352 - learning\_rate: 6.3096e-08  
Epoch 18/100  
31/31 ————— 0s 5ms/step - loss: 20.0785 - learning\_rate: 7.0795e-08  
Epoch 19/100  
31/31 ————— 0s 5ms/step - loss: 19.5876 - learning\_rate: 7.9433e-08  
Epoch 20/100  
31/31 ————— 0s 5ms/step - loss: 19.0520 - learning\_rate: 8.9125e-08  
Epoch 21/100  
31/31 ————— 0s 5ms/step - loss: 18.4632 - learning\_rate: 1.0000e-07  
Epoch 22/100  
31/31 ————— 0s 5ms/step - loss: 17.8136 - learning\_rate: 1.1220e-07  
Epoch 23/100  
31/31 ————— 0s 5ms/step - loss: 17.0986 - learning\_rate: 1.2589e-07  
Epoch 24/100  
31/31 ————— 0s 5ms/step - loss: 16.3107 - learning\_rate: 1.4125e-07  
Epoch 25/100  
31/31 ————— 0s 5ms/step - loss: 15.4398 - learning\_rate: 1.5849e-07  
Epoch 26/100  
31/31 ————— 0s 5ms/step - loss: 14.4814 - learning\_rate: 1.7783e-07  
Epoch 27/100  
31/31 ————— 0s 5ms/step - loss: 13.4190 - learning\_rate: 1.9953e-07  
Epoch 28/100  
31/31 ————— 0s 5ms/step - loss: 12.2325 - learning\_rate: 2.2387e-07  
Epoch 29/100  
31/31 ————— 0s 5ms/step - loss: 10.9616 - learning\_rate: 2.5119e-07  
Epoch 30/100  
31/31 ————— 0s 5ms/step - loss: 9.7130 - learning\_rate: 2.8184e-07  
Epoch 31/100  
31/31 ————— 0s 5ms/step - loss: 8.6223 - learning\_rate: 3.1623e-07  
Epoch 32/100  
31/31 ————— 0s 5ms/step - loss: 7.7907 - learning\_rate: 3.5481e-07  
Epoch 33/100  
31/31 ————— 0s 5ms/step - loss: 7.2619 - learning\_rate: 3.9811e-07  
Epoch 34/100  
31/31 ————— 0s 6ms/step - loss: 6.9522 - learning\_rate: 4.4668e-07  
Epoch 35/100  
31/31 ————— 0s 5ms/step - loss: 6.7418 - learning\_rate: 5.0119e-07  
Epoch 36/100  
31/31 ————— 0s 5ms/step - loss: 6.5812 - learning\_rate: 5.6234e-07  
Epoch 37/100  
31/31 ————— 0s 6ms/step - loss: 6.4447 - learning\_rate: 6.3096e-07  
Epoch 38/100  
31/31 ————— 0s 5ms/step - loss: 6.3066 - learning\_rate: 7.0795e-07  
Epoch 39/100  
31/31 ————— 0s 5ms/step - loss: 6.1649 - learning\_rate: 7.9433e-07

Epoch 40/100  
31/31  0s 5ms/step - loss: 6.0240 - learning\_rate: 8.9125e-07  
Epoch 41/100  
31/31  0s 5ms/step - loss: 5.8996 - learning\_rate: 1.0000e-06  
Epoch 42/100  
31/31  0s 5ms/step - loss: 5.7906 - learning\_rate: 1.1220e-06  
Epoch 43/100  
31/31  0s 5ms/step - loss: 5.6856 - learning\_rate: 1.2589e-06  
Epoch 44/100  
31/31  0s 5ms/step - loss: 5.5861 - learning\_rate: 1.4125e-06  
Epoch 45/100  
31/31  0s 5ms/step - loss: 5.4906 - learning\_rate: 1.5849e-06  
Epoch 46/100  
31/31  0s 5ms/step - loss: 5.3873 - learning\_rate: 1.7783e-06  
Epoch 47/100  
31/31  0s 5ms/step - loss: 5.2881 - learning\_rate: 1.9953e-06  
Epoch 48/100  
31/31  0s 5ms/step - loss: 5.2029 - learning\_rate: 2.2387e-06  
Epoch 49/100  
31/31  0s 5ms/step - loss: 5.1593 - learning\_rate: 2.5119e-06  
Epoch 50/100  
31/31  0s 5ms/step - loss: 5.1459 - learning\_rate: 2.8184e-06  
Epoch 51/100  
31/31  0s 5ms/step - loss: 5.1207 - learning\_rate: 3.1623e-06  
Epoch 52/100  
31/31  0s 6ms/step - loss: 5.0978 - learning\_rate: 3.5481e-06  
Epoch 53/100  
31/31  0s 6ms/step - loss: 5.1220 - learning\_rate: 3.9811e-06  
Epoch 54/100  
31/31  0s 6ms/step - loss: 5.1379 - learning\_rate: 4.4668e-06  
Epoch 55/100  
31/31  0s 6ms/step - loss: 4.9924 - learning\_rate: 5.0119e-06  
Epoch 56/100  
31/31  0s 6ms/step - loss: 4.7892 - learning\_rate: 5.6234e-06  
Epoch 57/100  
31/31  0s 5ms/step - loss: 4.7842 - learning\_rate: 6.3096e-06  
Epoch 58/100  
31/31  0s 6ms/step - loss: 4.6106 - learning\_rate: 7.0795e-06  
Epoch 59/100  
31/31  0s 6ms/step - loss: 4.6254 - learning\_rate: 7.9433e-06  
Epoch 60/100  
31/31  0s 6ms/step - loss: 4.7063 - learning\_rate: 8.9125e-06  
Epoch 61/100  
31/31  0s 5ms/step - loss: 4.6146 - learning\_rate: 1.0000e-05  
Epoch 62/100  
31/31  0s 5ms/step - loss: 4.5576 - learning\_rate: 1.1220e-05  
Epoch 63/100  
31/31  0s 5ms/step - loss: 4.6942 - learning\_rate: 1.2589e-05  
Epoch 64/100  
31/31  0s 5ms/step - loss: 4.7850 - learning\_rate: 1.4125e-05  
Epoch 65/100  
31/31  0s 5ms/step - loss: 4.6712 - learning\_rate: 1.5849e-05  
Epoch 66/100  
31/31  0s 5ms/step - loss: 4.8722 - learning\_rate: 1.7783e-05  
Epoch 67/100  
31/31  0s 5ms/step - loss: 5.0902 - learning\_rate: 1.9953e-05  
Epoch 68/100  
31/31  0s 6ms/step - loss: 5.1375 - learning\_rate: 2.2387e-05  
Epoch 69/100  
31/31  0s 6ms/step - loss: 5.3312 - learning\_rate: 2.5119e-05  
Epoch 70/100  
31/31  0s 5ms/step - loss: 5.0312 - learning\_rate: 2.8184e-05  
Epoch 71/100  
31/31  0s 5ms/step - loss: 5.6808 - learning\_rate: 3.1623e-05  
Epoch 72/100  
31/31  0s 6ms/step - loss: 5.7493 - learning\_rate: 3.5481e-05  
Epoch 73/100  
31/31  0s 6ms/step - loss: 5.5139 - learning\_rate: 3.9811e-05  
Epoch 74/100  
31/31  0s 6ms/step - loss: 5.6715 - learning\_rate: 4.4668e-05  
Epoch 75/100  
31/31  0s 6ms/step - loss: 6.2505 - learning\_rate: 5.0119e-05  
Epoch 76/100  
31/31  0s 6ms/step - loss: 6.0349 - learning\_rate: 5.6234e-05  
Epoch 77/100  
31/31  0s 6ms/step - loss: 6.6410 - learning\_rate: 6.3096e-05  
Epoch 78/100  
31/31  0s 5ms/step - loss: 9.0754 - learning\_rate: 7.0795e-05

```

Epoch 79/100
31/31 ————— 0s 5ms/step - loss: 7.4770 - learning_rate: 7.9433e-05
Epoch 80/100
31/31 ————— 0s 5ms/step - loss: 5.3599 - learning_rate: 8.9125e-05
Epoch 81/100
31/31 ————— 0s 5ms/step - loss: 7.5421 - learning_rate: 1.0000e-04
Epoch 82/100
31/31 ————— 0s 6ms/step - loss: 7.5513 - learning_rate: 1.1220e-04
Epoch 83/100
31/31 ————— 0s 6ms/step - loss: 9.0821 - learning_rate: 1.2589e-04
Epoch 84/100
31/31 ————— 0s 6ms/step - loss: 6.9848 - learning_rate: 1.4125e-04
Epoch 85/100
31/31 ————— 0s 6ms/step - loss: 7.7550 - learning_rate: 1.5849e-04
Epoch 86/100
31/31 ————— 0s 5ms/step - loss: 6.2030 - learning_rate: 1.7783e-04
Epoch 87/100
31/31 ————— 0s 5ms/step - loss: 7.5104 - learning_rate: 1.9953e-04
Epoch 88/100
31/31 ————— 0s 5ms/step - loss: 7.6245 - learning_rate: 2.2387e-04
Epoch 89/100
31/31 ————— 0s 5ms/step - loss: 7.9235 - learning_rate: 2.5119e-04
Epoch 90/100
31/31 ————— 0s 5ms/step - loss: 6.4471 - learning_rate: 2.8184e-04
Epoch 91/100
31/31 ————— 0s 6ms/step - loss: 8.5426 - learning_rate: 3.1623e-04
Epoch 92/100
31/31 ————— 0s 5ms/step - loss: 6.5610 - learning_rate: 3.5481e-04
Epoch 93/100
31/31 ————— 0s 5ms/step - loss: 8.1545 - learning_rate: 3.9811e-04
Epoch 94/100
31/31 ————— 0s 6ms/step - loss: 11.3164 - learning_rate: 4.4668e-04
Epoch 95/100
31/31 ————— 0s 5ms/step - loss: 6.8449 - learning_rate: 5.0119e-04
Epoch 96/100
31/31 ————— 0s 5ms/step - loss: 8.0550 - learning_rate: 5.6234e-04
Epoch 97/100
31/31 ————— 0s 5ms/step - loss: 10.4993 - learning_rate: 6.3096e-04
Epoch 98/100
31/31 ————— 0s 5ms/step - loss: 9.3931 - learning_rate: 7.0795e-04
Epoch 99/100
31/31 ————— 0s 5ms/step - loss: 10.5945 - learning_rate: 7.9433e-04
Epoch 100/100
31/31 ————— 0s 5ms/step - loss: 10.4864 - learning_rate: 8.9125e-04

```

```

In [10]: # Define the Learning rate array
lrs = 1e-8 * (10 ** (np.arange(100) / 20))

# Set the figure size
plt.figure(figsize=(10, 6))

# Set the grid
plt.grid(True)

# Plot the loss in log scale
plt.semilogx(lrs, history.history["loss"])

# Increase the tickmarks size
plt.tick_params('both', length=10, width=1, which='both')


# Set the plot boundaries
plt.axis([1e-8, 1e-3, 0, 30])

```

```

Out[10]: (1e-08, 0.001, 0.0, 30.0)

```



## Train the Model

You can then proceed to train the model with your chosen learning rate.

*Tip: When experimenting and you find yourself running different iterations of a model, you may want to use the `clear_session()` method to declutter memory used by Keras. This is added in the first line below.*

```
In [11]: # Reset states generated by Keras
tf.keras.backend.clear_session()

# Build the model
model = tf.keras.models.Sequential([
    tf.keras.Input(shape=(window_size, 1)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 100.0)
])

# Set the Learning rate
learning_rate = 2e-6

# Set the optimizer
optimizer = tf.keras.optimizers.SGD(learning_rate=learning_rate, momentum=0.9)

# Set the training parameters
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])

# Train the model
history = model.fit(dataset, epochs=100)
```



Epoch 1/100  
31/31 ————— 2s 6ms/step - loss: 66.4583 - mae: 66.9578  
Epoch 2/100  
31/31 ————— 0s 6ms/step - loss: 13.0632 - mae: 13.5602  
Epoch 3/100  
31/31 ————— 0s 6ms/step - loss: 9.1585 - mae: 9.6362  
Epoch 4/100  
31/31 ————— 0s 6ms/step - loss: 8.4469 - mae: 8.9343  
Epoch 5/100  
31/31 ————— 0s 6ms/step - loss: 8.0986 - mae: 8.5859  
Epoch 6/100  
31/31 ————— 0s 6ms/step - loss: 7.7979 - mae: 8.2862  
Epoch 7/100  
31/31 ————— 0s 6ms/step - loss: 7.4471 - mae: 7.9360  
Epoch 8/100  
31/31 ————— 0s 6ms/step - loss: 6.9718 - mae: 7.4576  
Epoch 9/100  
31/31 ————— 0s 6ms/step - loss: 6.6138 - mae: 7.0946  
Epoch 10/100  
31/31 ————— 0s 6ms/step - loss: 6.4142 - mae: 6.8958  
Epoch 11/100  
31/31 ————— 0s 6ms/step - loss: 6.2424 - mae: 6.7223  
Epoch 12/100  
31/31 ————— 0s 6ms/step - loss: 6.1135 - mae: 6.5937  
Epoch 13/100  
31/31 ————— 0s 6ms/step - loss: 6.0016 - mae: 6.4827  
Epoch 14/100  
31/31 ————— 0s 6ms/step - loss: 5.8965 - mae: 6.3775  
Epoch 15/100  
31/31 ————— 0s 6ms/step - loss: 5.8003 - mae: 6.2805  
Epoch 16/100  
31/31 ————— 0s 6ms/step - loss: 5.7236 - mae: 6.2035  
Epoch 17/100  
31/31 ————— 0s 6ms/step - loss: 5.6583 - mae: 6.1391  
Epoch 18/100  
31/31 ————— 0s 6ms/step - loss: 5.5891 - mae: 6.0658  
Epoch 19/100  
31/31 ————— 0s 6ms/step - loss: 5.5209 - mae: 6.0012  
Epoch 20/100  
31/31 ————— 0s 6ms/step - loss: 5.4769 - mae: 5.9569  
Epoch 21/100  
31/31 ————— 0s 6ms/step - loss: 5.4207 - mae: 5.9000  
Epoch 22/100  
31/31 ————— 0s 6ms/step - loss: 5.4302 - mae: 5.9100  
Epoch 23/100  
31/31 ————— 0s 6ms/step - loss: 5.4410 - mae: 5.9174  
Epoch 24/100  
31/31 ————— 0s 6ms/step - loss: 5.3457 - mae: 5.8226  
Epoch 25/100  
31/31 ————— 0s 6ms/step - loss: 5.2259 - mae: 5.7030  
Epoch 26/100  
31/31 ————— 0s 6ms/step - loss: 5.2093 - mae: 5.6842  
Epoch 27/100  
31/31 ————— 0s 6ms/step - loss: 5.1578 - mae: 5.6312  
Epoch 28/100  
31/31 ————— 0s 6ms/step - loss: 5.0689 - mae: 5.5448  
Epoch 29/100  
31/31 ————— 0s 6ms/step - loss: 5.0629 - mae: 5.5380  
Epoch 30/100  
31/31 ————— 0s 6ms/step - loss: 5.0776 - mae: 5.5513  
Epoch 31/100  
31/31 ————— 0s 6ms/step - loss: 5.0193 - mae: 5.4938  
Epoch 32/100  
31/31 ————— 0s 5ms/step - loss: 4.9592 - mae: 5.4346  
Epoch 33/100  
31/31 ————— 0s 5ms/step - loss: 4.9672 - mae: 5.4412  
Epoch 34/100  
31/31 ————— 0s 5ms/step - loss: 4.9442 - mae: 5.4185  
Epoch 35/100  
31/31 ————— 0s 5ms/step - loss: 4.8918 - mae: 5.3673  
Epoch 36/100  
31/31 ————— 0s 5ms/step - loss: 4.9072 - mae: 5.3808  
Epoch 37/100  
31/31 ————— 0s 5ms/step - loss: 4.8872 - mae: 5.3614  
Epoch 38/100  
31/31 ————— 0s 5ms/step - loss: 4.8553 - mae: 5.3288  
Epoch 39/100  
31/31 ————— 0s 6ms/step - loss: 4.8453 - mae: 5.3191

Epoch 40/100  
31/31 ————— 0s 6ms/step - loss: 4.8168 - mae: 5.2904  
Epoch 41/100  
31/31 ————— 0s 6ms/step - loss: 4.8075 - mae: 5.2803  
Epoch 42/100  
31/31 ————— 0s 6ms/step - loss: 4.7873 - mae: 5.2610  
Epoch 43/100  
31/31 ————— 0s 6ms/step - loss: 4.7738 - mae: 5.2472  
Epoch 44/100  
31/31 ————— 0s 6ms/step - loss: 4.7603 - mae: 5.2334  
Epoch 45/100  
31/31 ————— 0s 6ms/step - loss: 4.7484 - mae: 5.2206  
Epoch 46/100  
31/31 ————— 0s 6ms/step - loss: 4.7369 - mae: 5.2088  
Epoch 47/100  
31/31 ————— 0s 6ms/step - loss: 4.7270 - mae: 5.1978  
Epoch 48/100  
31/31 ————— 0s 6ms/step - loss: 4.7168 - mae: 5.1871  
Epoch 49/100  
31/31 ————— 0s 6ms/step - loss: 4.7077 - mae: 5.1774  
Epoch 50/100  
31/31 ————— 0s 6ms/step - loss: 4.6993 - mae: 5.1686  
Epoch 51/100  
31/31 ————— 0s 6ms/step - loss: 4.6913 - mae: 5.1611  
Epoch 52/100  
31/31 ————— 0s 6ms/step - loss: 4.6828 - mae: 5.1531  
Epoch 53/100  
31/31 ————— 0s 6ms/step - loss: 4.6748 - mae: 5.1453  
Epoch 54/100  
31/31 ————— 0s 6ms/step - loss: 4.6677 - mae: 5.1388  
Epoch 55/100  
31/31 ————— 0s 6ms/step - loss: 4.6599 - mae: 5.1317  
Epoch 56/100  
31/31 ————— 0s 6ms/step - loss: 4.6535 - mae: 5.1257  
Epoch 57/100  
31/31 ————— 0s 6ms/step - loss: 4.6466 - mae: 5.1190  
Epoch 58/100  
31/31 ————— 0s 6ms/step - loss: 4.6404 - mae: 5.1131  
Epoch 59/100  
31/31 ————— 0s 6ms/step - loss: 4.6342 - mae: 5.1069  
Epoch 60/100  
31/31 ————— 0s 6ms/step - loss: 4.6285 - mae: 5.1013  
Epoch 61/100  
31/31 ————— 0s 6ms/step - loss: 4.6224 - mae: 5.0951  
Epoch 62/100  
31/31 ————— 0s 6ms/step - loss: 4.6165 - mae: 5.0889  
Epoch 63/100  
31/31 ————— 0s 6ms/step - loss: 4.6109 - mae: 5.0833  
Epoch 64/100  
31/31 ————— 0s 6ms/step - loss: 4.6044 - mae: 5.0767  
Epoch 65/100  
31/31 ————— 0s 6ms/step - loss: 4.5984 - mae: 5.0706  
Epoch 66/100  
31/31 ————— 0s 6ms/step - loss: 4.5917 - mae: 5.0639  
Epoch 67/100  
31/31 ————— 0s 6ms/step - loss: 4.5849 - mae: 5.0572  
Epoch 68/100  
31/31 ————— 0s 6ms/step - loss: 4.5802 - mae: 5.0519  
Epoch 69/100  
31/31 ————— 0s 6ms/step - loss: 4.5744 - mae: 5.0462  
Epoch 70/100  
31/31 ————— 0s 6ms/step - loss: 4.5687 - mae: 5.0402  
Epoch 71/100  
31/31 ————— 0s 6ms/step - loss: 4.5623 - mae: 5.0338  
Epoch 72/100  
31/31 ————— 0s 6ms/step - loss: 4.5565 - mae: 5.0275  
Epoch 73/100  
31/31 ————— 0s 6ms/step - loss: 4.5502 - mae: 5.0212  
Epoch 74/100  
31/31 ————— 0s 6ms/step - loss: 4.5441 - mae: 5.0149  
Epoch 75/100  
31/31 ————— 0s 6ms/step - loss: 4.5389 - mae: 5.0093  
Epoch 76/100  
31/31 ————— 0s 6ms/step - loss: 4.5349 - mae: 5.0053  
Epoch 77/100  
31/31 ————— 0s 6ms/step - loss: 4.5279 - mae: 4.9984  
Epoch 78/100  
31/31 ————— 0s 6ms/step - loss: 4.5234 - mae: 4.9940

```

Epoch 79/100
31/31 ————— 0s 6ms/step - loss: 4.5190 - mae: 4.9898
Epoch 80/100
31/31 ————— 0s 6ms/step - loss: 4.5141 - mae: 4.9848
Epoch 81/100
31/31 ————— 0s 6ms/step - loss: 4.5092 - mae: 4.9799
Epoch 82/100
31/31 ————— 0s 6ms/step - loss: 4.5056 - mae: 4.9766
Epoch 83/100
31/31 ————— 0s 6ms/step - loss: 4.5009 - mae: 4.9720
Epoch 84/100
31/31 ————— 0s 6ms/step - loss: 4.4963 - mae: 4.9675
Epoch 85/100
31/31 ————— 0s 6ms/step - loss: 4.4951 - mae: 4.9660
Epoch 86/100
31/31 ————— 0s 6ms/step - loss: 4.4903 - mae: 4.9613
Epoch 87/100
31/31 ————— 0s 6ms/step - loss: 4.4872 - mae: 4.9578
Epoch 88/100
31/31 ————— 0s 6ms/step - loss: 4.4835 - mae: 4.9538
Epoch 89/100
31/31 ————— 0s 6ms/step - loss: 4.4831 - mae: 4.9543
Epoch 90/100
31/31 ————— 0s 6ms/step - loss: 4.4776 - mae: 4.9493
Epoch 91/100
31/31 ————— 0s 6ms/step - loss: 4.4737 - mae: 4.9466
Epoch 92/100
31/31 ————— 0s 6ms/step - loss: 4.4731 - mae: 4.9465
Epoch 93/100
31/31 ————— 0s 6ms/step - loss: 4.4712 - mae: 4.9454
Epoch 94/100
31/31 ————— 0s 5ms/step - loss: 4.4694 - mae: 4.9441
Epoch 95/100
31/31 ————— 0s 6ms/step - loss: 4.4654 - mae: 4.9402
Epoch 96/100
31/31 ————— 0s 5ms/step - loss: 4.4672 - mae: 4.9421
Epoch 97/100
31/31 ————— 0s 6ms/step - loss: 4.4658 - mae: 4.9409
Epoch 98/100
31/31 ————— 0s 6ms/step - loss: 4.4634 - mae: 4.9385
Epoch 99/100
31/31 ————— 0s 6ms/step - loss: 4.4611 - mae: 4.9362
Epoch 100/100
31/31 ————— 0s 6ms/step - loss: 4.4599 - mae: 4.9349

```

## Model Prediction

You will then generate batches of windows to generate predictions that align with the validation set.

```

In [12]: def model_forecast(model, series, window_size, batch_size):
          """Uses an input model to generate predictions on data windows

          Args:
            model (TF Keras Model) - model that accepts data windows
            series (array of float) - contains the values of the time series
            window_size (int) - the number of time steps to include in the window
            batch_size (int) - the batch size

          Returns:
            forecast (numpy array) - array containing predictions
          """

          # Add an axis for the feature dimension of RNN Layers
          series = tf.expand_dims(series, axis=-1)

          # Generate a TF Dataset from the series values
          dataset = tf.data.Dataset.from_tensor_slices(series)

          # Window the data but only take those with the specified size
          dataset = dataset.window(window_size, shift=1, drop_remainder=True)

          # Flatten the windows by putting its elements in a single batch
          dataset = dataset.flat_map(lambda w: w.batch(window_size))

          # Create batches of windows
          dataset = dataset.batch(batch_size).prefetch(1)

```

```
# Get predictions on the entire dataset
forecast = model.predict(dataset, verbose=0)


return forecast
```

```
In [13]: # Reduce the original series
forecast_series = series[split_time-window_size:-1]

# Use helper function to generate predictions
forecast = model_forecast(model, forecast_series, window_size, batch_size)

# Drop single dimensional axis
results = forecast.squeeze()

# Plot the results
plot_series(time_valid, (x_valid, results))
```



You can then generate the metrics to evaluate the model's performance.

```
In [14]: # Compute the MSE and MAE
print(tf.keras.metrics.mse(x_valid, results).numpy())
print(tf.keras.metrics.mae(x_valid, results).numpy())
```

```
67.75228
5.7619185
```

## Wrap Up

This concludes this short exercise on using LSTMs for time series forecasting. Next week, you will build upon this and add convolutions. Then, you will start to move away from synthetic data and use real-world datasets. See you there!

## Optional: Including a Validation Set while Training

Back in the first course of this specialization, you saw how you can also monitor the performance of your model against a validation set while training. You can also do that for this lab.

First, you need to generate a `val_set` which are data windows and labels that your model can accept. You can simply reuse the `windowed_dataset` function for that and you can pass in the `x_valid` points to generate the windows.

```
In [15]: # Generate data windows of the validation set
val_set = windowed_dataset(x_valid, window_size, batch_size, shuffle_buffer_size)
```

You can then do the same training as before but pass in the `val_set` to the `validation_data` parameter of the `fit()` method.

```
In [16]: # Reset states generated by Keras
tf.keras.backend.clear_session()

# Build the model
model = tf.keras.models.Sequential([
    tf.keras.Input(shape=(window_size, 1)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(1),
    tf.keras.layers.Lambda(lambda x: x * 100.0)
])

# Set the Learning rate
learning_rate = 2e-6























# Set the optimizer
optimizer = tf.keras.optimizers.SGD(learning_rate=learning_rate, momentum=0.9)

# Set the training parameters
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])

# Train the model
history = model.fit(dataset, epochs=100, validation_data=val_set)
```

Epoch 1/100  
31/31 ————— 2s 20ms/step - loss: 32.7407 - mae: 33.2375 - val\_loss: 22.3182 - val\_mae: 22.8143  
Epoch 2/100  
31/31 ————— 0s 7ms/step - loss: 12.1673 - mae: 12.6610 - val\_loss: 17.5748 - val\_mae: 18.0723  
Epoch 3/100  
31/31 ————— 0s 7ms/step - loss: 7.5763 - mae: 8.0522 - val\_loss: 12.2577 - val\_mae: 12.7495  
Epoch 4/100  
31/31 ————— 0s 7ms/step - loss: 6.6635 - mae: 7.1468 - val\_loss: 10.9568 - val\_mae: 11.4456  
Epoch 5/100  
31/31 ————— 0s 7ms/step - loss: 6.4514 - mae: 6.9334 - val\_loss: 10.5222 - val\_mae: 11.0129  
Epoch 6/100  
31/31 ————— 0s 7ms/step - loss: 6.2836 - mae: 6.7668 - val\_loss: 10.1590 - val\_mae: 10.6476  
Epoch 7/100  
31/31 ————— 0s 7ms/step - loss: 6.1421 - mae: 6.6274 - val\_loss: 9.8933 - val\_mae: 10.3809  
Epoch 8/100  
31/31 ————— 0s 7ms/step - loss: 6.0093 - mae: 6.4942 - val\_loss: 9.6813 - val\_mae: 10.1698  
Epoch 9/100  
31/31 ————— 0s 7ms/step - loss: 5.8863 - mae: 6.3704 - val\_loss: 9.4973 - val\_mae: 9.9870  
Epoch 10/100  
31/31 ————— 0s 7ms/step - loss: 5.7738 - mae: 6.2579 - val\_loss: 9.2777 - val\_mae: 9.7662  
Epoch 11/100  
31/31 ————— 0s 7ms/step - loss: 5.6705 - mae: 6.1548 - val\_loss: 9.0967 - val\_mae: 9.5842  
Epoch 12/100  
31/31 ————— 0s 7ms/step - loss: 5.5715 - mae: 6.0554 - val\_loss: 8.9124 - val\_mae: 9.3985  
Epoch 13/100  
31/31 ————— 0s 7ms/step - loss: 5.4816 - mae: 5.9662 - val\_loss: 8.7294 - val\_mae: 9.2144  
Epoch 14/100  
31/31 ————— 0s 7ms/step - loss: 5.4034 - mae: 5.8877 - val\_loss: 8.5282 - val\_mae: 9.0127  
Epoch 15/100  
31/31 ————— 0s 7ms/step - loss: 5.3326 - mae: 5.8144 - val\_loss: 8.3353 - val\_mae: 8.8217  
Epoch 16/100  
31/31 ————— 0s 7ms/step - loss: 5.2656 - mae: 5.7466 - val\_loss: 8.1559 - val\_mae: 8.6416  
Epoch 17/100  
31/31 ————— 0s 7ms/step - loss: 5.2053 - mae: 5.6837 - val\_loss: 7.9849 - val\_mae: 8.4695  
Epoch 18/100  
31/31 ————— 0s 7ms/step - loss: 5.1496 - mae: 5.6280 - val\_loss: 7.8284 - val\_mae: 8.3121  
Epoch 19/100  
31/31 ————— 0s 7ms/step - loss: 5.0995 - mae: 5.5794 - val\_loss: 7.6667 - val\_mae: 8.1507  
Epoch 20/100  
31/31 ————— 0s 8ms/step - loss: 5.0525 - mae: 5.5340 - val\_loss: 7.5235 - val\_mae: 8.0081  
Epoch 21/100  
31/31 ————— 0s 7ms/step - loss: 5.0088 - mae: 5.4902 - val\_loss: 7.3718 - val\_mae: 7.8558  
Epoch 22/100  
31/31 ————— 0s 7ms/step - loss: 4.9697 - mae: 5.4509 - val\_loss: 7.2433 - val\_mae: 7.7270  
Epoch 23/100  
31/31 ————— 0s 7ms/step - loss: 4.9356 - mae: 5.4172 - val\_loss: 7.1243 - val\_mae: 7.6078  
Epoch 24/100  
31/31 ————— 0s 7ms/step - loss: 4.9049 - mae: 5.3861 - val\_loss: 6.9928 - val\_mae: 7.4759  
Epoch 25/100  
31/31 ————— 0s 7ms/step - loss: 4.8762 - mae: 5.3568 - val\_loss: 6.8856 - val\_mae: 7.3666  
Epoch 26/100  
31/31 ————— 0s 7ms/step - loss: 4.8490 - mae: 5.3294 - val\_loss: 6.7913 - val\_mae: 7.2710  
Epoch 27/100  
31/31 ————— 0s 7ms/step - loss: 4.8234 - mae: 5.3043 - val\_loss: 6.6893 - val\_mae: 7.1695  
Epoch 28/100  
31/31 ————— 0s 7ms/step - loss: 4.8007 - mae: 5.2826 - val\_loss: 6.6051 - val\_mae: 7.0867  
Epoch 29/100  
31/31 ————— 0s 7ms/step - loss: 4.7772 - mae: 5.2603 - val\_loss: 6.5213 - val\_mae: 7.0046  
Epoch 30/100  
31/31 ————— 0s 7ms/step - loss: 4.7569 - mae: 5.2409 - val\_loss: 6.4488 - val\_mae: 6.9322  
Epoch 31/100  
31/31 ————— 0s 7ms/step - loss: 4.7361 - mae: 5.2207 - val\_loss: 6.3685 - val\_mae: 6.8514  
Epoch 32/100  
31/31 ————— 0s 7ms/step - loss: 4.7185 - mae: 5.2031 - val\_loss: 6.2921 - val\_mae: 6.7734  
Epoch 33/100  
31/31 ————— 0s 7ms/step - loss: 4.7022 - mae: 5.1866 - val\_loss: 6.2262 - val\_mae: 6.7063  
Epoch 34/100  
31/31 ————— 0s 7ms/step - loss: 4.6849 - mae: 5.1693 - val\_loss: 6.1569 - val\_mae: 6.6366  
Epoch 35/100  
31/31 ————— 0s 7ms/step - loss: 4.6700 - mae: 5.1546 - val\_loss: 6.0915 - val\_mae: 6.5731  
Epoch 36/100  
31/31 ————— 0s 7ms/step - loss: 4.6556 - mae: 5.1402 - val\_loss: 6.0228 - val\_mae: 6.5063  
Epoch 37/100  
31/31 ————— 0s 7ms/step - loss: 4.6436 - mae: 5.1277 - val\_loss: 5.9600 - val\_mae: 6.4455  
Epoch 38/100  
31/31 ————— 0s 7ms/step - loss: 4.6337 - mae: 5.1170 - val\_loss: 5.8997 - val\_mae: 6.3855  
Epoch 39/100  
31/31 ————— 0s 7ms/step - loss: 4.6229 - mae: 5.1052 - val\_loss: 5.8415 - val\_mae: 6.3266

Epoch 40/100  
31/31 — 0s 7ms/step - loss: 4.6148 - mae: 5.0966 - val\_loss: 5.7922 - val\_mae: 6.2759  
Epoch 41/100  
31/31 — 0s 7ms/step - loss: 4.6074 - mae: 5.0898 - val\_loss: 5.7488 - val\_mae: 6.2310  
Epoch 42/100  
31/31 — 0s 7ms/step - loss: 4.6022 - mae: 5.0851 - val\_loss: 5.7136 - val\_mae: 6.1948  
Epoch 43/100  
31/31 — 0s 7ms/step - loss: 4.5977 - mae: 5.0801 - val\_loss: 5.6800 - val\_mae: 6.1602  
Epoch 44/100  
31/31 — 0s 7ms/step - loss: 4.5970 - mae: 5.0785 - val\_loss: 5.6620 - val\_mae: 6.1420  
Epoch 45/100  
31/31 — 0s 7ms/step - loss: 4.5955 - mae: 5.0765 - val\_loss: 5.6343 - val\_mae: 6.1135  
Epoch 46/100  
31/31 — 0s 7ms/step - loss: 4.5952 - mae: 5.0757 - val\_loss: 5.6149 - val\_mae: 6.0935  
Epoch 47/100  
31/31 — 0s 7ms/step - loss: 4.5916 - mae: 5.0716 - val\_loss: 5.5969 - val\_mae: 6.0750  
Epoch 48/100  
31/31 — 0s 7ms/step - loss: 4.5879 - mae: 5.0671 - val\_loss: 5.5776 - val\_mae: 6.0554  
Epoch 49/100  
31/31 — 0s 7ms/step - loss: 4.5851 - mae: 5.0639 - val\_loss: 5.5592 - val\_mae: 6.0370  
Epoch 50/100  
31/31 — 0s 7ms/step - loss: 4.5838 - mae: 5.0622 - val\_loss: 5.5416 - val\_mae: 6.0193  
Epoch 51/100  
31/31 — 0s 7ms/step - loss: 4.5824 - mae: 5.0601 - val\_loss: 5.5252 - val\_mae: 6.0026  
Epoch 52/100  
31/31 — 0s 7ms/step - loss: 4.5796 - mae: 5.0571 - val\_loss: 5.5108 - val\_mae: 5.9879  
Epoch 53/100  
31/31 — 0s 7ms/step - loss: 4.5776 - mae: 5.0546 - val\_loss: 5.4933 - val\_mae: 5.9698  
Epoch 54/100  
31/31 — 0s 7ms/step - loss: 4.5752 - mae: 5.0518 - val\_loss: 5.4780 - val\_mae: 5.9542  
Epoch 55/100  
31/31 — 0s 7ms/step - loss: 4.5717 - mae: 5.0479 - val\_loss: 5.4667 - val\_mae: 5.9428  
Epoch 56/100  
31/31 — 0s 7ms/step - loss: 4.5678 - mae: 5.0438 - val\_loss: 5.4545 - val\_mae: 5.9305  
Epoch 57/100  
31/31 — 0s 7ms/step - loss: 4.5671 - mae: 5.0427 - val\_loss: 5.4443 - val\_mae: 5.9202  
Epoch 58/100  
31/31 — 0s 7ms/step - loss: 4.5653 - mae: 5.0407 - val\_loss: 5.4292 - val\_mae: 5.9049  
Epoch 59/100  
31/31 — 0s 7ms/step - loss: 4.5665 - mae: 5.0415 - val\_loss: 5.4206 - val\_mae: 5.8963  
Epoch 60/100  
31/31 — 0s 7ms/step - loss: 4.5647 - mae: 5.0396 - val\_loss: 5.4108 - val\_mae: 5.8864  
Epoch 61/100  
31/31 — 0s 7ms/step - loss: 4.5623 - mae: 5.0372 - val\_loss: 5.4002 - val\_mae: 5.8757  
Epoch 62/100  
31/31 — 0s 7ms/step - loss: 4.5585 - mae: 5.0334 - val\_loss: 5.3895 - val\_mae: 5.8649  
Epoch 63/100  
31/31 — 0s 7ms/step - loss: 4.5558 - mae: 5.0307 - val\_loss: 5.3799 - val\_mae: 5.8550  
Epoch 64/100  
31/31 — 0s 7ms/step - loss: 4.5536 - mae: 5.0282 - val\_loss: 5.3691 - val\_mae: 5.8438  
Epoch 65/100  
31/31 — 0s 7ms/step - loss: 4.5507 - mae: 5.0254 - val\_loss: 5.3596 - val\_mae: 5.8338  
Epoch 66/100  
31/31 — 0s 7ms/step - loss: 4.5497 - mae: 5.0244 - val\_loss: 5.3521 - val\_mae: 5.8263  
Epoch 67/100  
31/31 — 0s 7ms/step - loss: 4.5500 - mae: 5.0247 - val\_loss: 5.3417 - val\_mae: 5.8159  
Epoch 68/100  
31/31 — 0s 7ms/step - loss: 4.5482 - mae: 5.0228 - val\_loss: 5.3317 - val\_mae: 5.8060  
Epoch 69/100  
31/31 — 0s 7ms/step - loss: 4.5449 - mae: 5.0193 - val\_loss: 5.3245 - val\_mae: 5.7986  
Epoch 70/100  
31/31 — 0s 7ms/step - loss: 4.5414 - mae: 5.0157 - val\_loss: 5.3187 - val\_mae: 5.7926  
Epoch 71/100  
31/31 — 0s 7ms/step - loss: 4.5403 - mae: 5.0148 - val\_loss: 5.3119 - val\_mae: 5.7857  
Epoch 72/100  
31/31 — 0s 7ms/step - loss: 4.5403 - mae: 5.0150 - val\_loss: 5.3008 - val\_mae: 5.7747  
Epoch 73/100  
31/31 — 0s 7ms/step - loss: 4.5405 - mae: 5.0151 - val\_loss: 5.2955 - val\_mae: 5.7692  
Epoch 74/100  
31/31 — 0s 7ms/step - loss: 4.5391 - mae: 5.0138 - val\_loss: 5.2897 - val\_mae: 5.7633  
Epoch 75/100  
31/31 — 0s 7ms/step - loss: 4.5378 - mae: 5.0126 - val\_loss: 5.2818 - val\_mae: 5.7551  
Epoch 76/100  
31/31 — 0s 7ms/step - loss: 4.5361 - mae: 5.0108 - val\_loss: 5.2743 - val\_mae: 5.7472  
Epoch 77/100  
31/31 — 0s 7ms/step - loss: 4.5325 - mae: 5.0071 - val\_loss: 5.2668 - val\_mae: 5.7394  
Epoch 78/100  
31/31 — 0s 7ms/step - loss: 4.5297 - mae: 5.0041 - val\_loss: 5.2608 - val\_mae: 5.7333

Epoch 79/100  
31/31  0s 7ms/step - loss: 4.5271 - mae: 5.0013 - val\_loss: 5.2526 - val\_mae: 5.7249  
Epoch 80/100  
31/31  0s 7ms/step - loss: 4.5252 - mae: 4.9993 - val\_loss: 5.2458 - val\_mae: 5.7178  
Epoch 81/100  
31/31  0s 7ms/step - loss: 4.5213 - mae: 4.9952 - val\_loss: 5.2408 - val\_mae: 5.7125  
Epoch 82/100  
31/31  0s 7ms/step - loss: 4.5189 - mae: 4.9929 - val\_loss: 5.2351 - val\_mae: 5.7066  
Epoch 83/100  
31/31  0s 8ms/step - loss: 4.5177 - mae: 4.9917 - val\_loss: 5.2313 - val\_mae: 5.7027  
Epoch 84/100  
31/31  0s 7ms/step - loss: 4.5144 - mae: 4.9884 - val\_loss: 5.2235 - val\_mae: 5.6947  
Epoch 85/100  
31/31  0s 7ms/step - loss: 4.5118 - mae: 4.9855 - val\_loss: 5.2161 - val\_mae: 5.6871  
Epoch 86/100  
31/31  0s 7ms/step - loss: 4.5109 - mae: 4.9846 - val\_loss: 5.2133 - val\_mae: 5.6841  
Epoch 87/100  
31/31  0s 7ms/step - loss: 4.5102 - mae: 4.9843 - val\_loss: 5.2095 - val\_mae: 5.6802  
Epoch 88/100  
31/31  0s 7ms/step - loss: 4.5089 - mae: 4.9830 - val\_loss: 5.2033 - val\_mae: 5.6738  
Epoch 89/100  
31/31  0s 7ms/step - loss: 4.5069 - mae: 4.9810 - val\_loss: 5.1995 - val\_mae: 5.6698  
Epoch 90/100  
31/31  0s 7ms/step - loss: 4.5044 - mae: 4.9786 - val\_loss: 5.1953 - val\_mae: 5.6656  
Epoch 91/100  
31/31  0s 7ms/step - loss: 4.5029 - mae: 4.9772 - val\_loss: 5.1908 - val\_mae: 5.6610  
Epoch 92/100  
31/31  0s 7ms/step - loss: 4.5031 - mae: 4.9775 - val\_loss: 5.1859 - val\_mae: 5.6562  
Epoch 93/100  
31/31  0s 7ms/step - loss: 4.5035 - mae: 4.9780 - val\_loss: 5.1819 - val\_mae: 5.6522  
Epoch 94/100  
31/31  0s 7ms/step - loss: 4.5026 - mae: 4.9771 - val\_loss: 5.1791 - val\_mae: 5.6493  
Epoch 95/100  
31/31  0s 7ms/step - loss: 4.4996 - mae: 4.9739 - val\_loss: 5.1737 - val\_mae: 5.6438  
Epoch 96/100  
31/31  0s 7ms/step - loss: 4.4980 - mae: 4.9723 - val\_loss: 5.1715 - val\_mae: 5.6416  
Epoch 97/100  
31/31  0s 7ms/step - loss: 4.4938 - mae: 4.9681 - val\_loss: 5.1681 - val\_mae: 5.6381  
Epoch 98/100  
31/31  0s 7ms/step - loss: 4.4916 - mae: 4.9658 - val\_loss: 5.1627 - val\_mae: 5.6326  
Epoch 99/100  
31/31  0s 7ms/step - loss: 4.4912 - mae: 4.9653 - val\_loss: 5.1595 - val\_mae: 5.6293  
Epoch 100/100  
31/31  0s 8ms/step - loss: 4.4899 - mae: 4.9640 - val\_loss: 5.1534 - val\_mae: 5.6234