Ungraded Lab: Predicting Sunspots with Neural Networks

At this point in the course, you should be able to explore different network architectures for forecasting. In the previous weeks, you've used DNNs, RNNs, and CNNs to build these different models. In the final practice lab for this course, you'll try one more configuration and that is a combination of all these types of networks: the data windows will pass through a convolution, followed by stacked LSTMs, followed by stacked dense layers. See if this improves results or you can just opt for simpler models.

Imports

```
In [1]: import tensorflow as tf
    import numpy as np
    import matplotlib.pyplot as plt
    import csv
```

Utilities

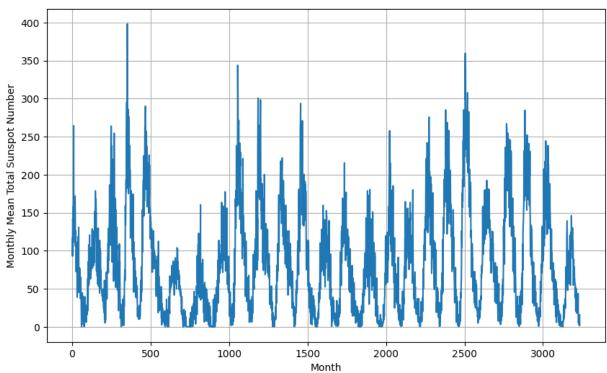
```
In [2]: def plot_series(x, y, format="-", start=0, end=None,
                        title=None, xlabel=None, ylabel=None, legend=None ):
            Visualizes time series data
              x (array of int) - contains values for the x-axis
              y (array of int or tuple of arrays) - contains the values for the y-axis
              format (string) - line style when plotting the graph
              start (int) - first time step to plot
              end (int) - last time step to plot
              title (string) - title of the plot
              xlabel (string) - label for the x-axis
              ylabel (string) - label for the y-axis
              legend (list of strings) - legend for the plot
            # Setup dimensions of the graph figure
            plt.figure(figsize=(10, 6))
            # Check if there are more than two series to plot
            if type(y) is tuple:
              # Loop over the y elements
              for y_curr in y:
                # Plot the x and current y values
                plt.plot(x[start:end], y_curr[start:end], format)
            else:
              \# Plot the x and y values
              plt.plot(x[start:end], y[start:end], format)
            # Label the x-axis
            plt.xlabel(xlabel)
            # Label the y-axis
            plt.ylabel(ylabel)
            # Set the Legend
            if legend:
              plt.legend(legend)
            # Set the title
            plt.title(title)
            # Overlay a grid on the graph
            plt.grid(True)
            # Draw the graph on screen
            plt.show()
```

Download and Preview the Dataset

```
In [3]: # Download the Dataset
!wget -nc https://storage.googleapis.com/tensorflow-1-public/course4/Sunspots.csv
```

File 'Sunspots.csv' already there; not retrieving.

```
In [4]: # Initialize lists
        time_step = []
        sunspots = []
        # Open CSV file
        with open('./Sunspots.csv') as csvfile:
          # Initialize reader
          reader = csv.reader(csvfile, delimiter=',')
          # Skip the first line
          next(reader)
          # Append row and sunspot number to lists
          for row in reader:
            time_step.append(int(row[0]))
            {\tt sunspots.append(float(row[2]))}
        # Convert lists to numpy arrays
        time = np.array(time_step)
        series = np.array(sunspots)
        # Preview the data
        plot_series(time, series, xlabel='Month', ylabel='Monthly Mean Total Sunspot Number')
```



Split the Dataset

```
In [5]: # Define the split time
    split_time = 3000

# Get the train set
    time_train = time[:split_time]
    x_train = series[:split_time]
# Get the validation set
```

```
time_valid = time[split_time:]
x_valid = series[split_time:]
```

Prepare Features and Labels

```
In [6]: def windowed_dataset(series, window_size, batch_size, shuffle_buffer):
                                                        """Generates dataset windows
                                                              series (array of float) - contains the values of the time series % \left( \frac{1}{2}\right) =\left( \frac{1}{2}\right) \left( \frac{1}{2}\right) 
                                                             window_size (int) - the number of time steps to include in the feature
                                                             batch_size (int) - the batch size
                                                             shuffle_buffer(int) - buffer size to use for the shuffle method
                                                            dataset (TF Dataset) - TF Dataset containing time windows
                                                     # Add an axis for the feature dimension of RNN layers
                                                     series = tf.expand_dims(series, axis=-1)
                                                     # Generate a TF Dataset from the series values
                                                     dataset = tf.data.Dataset.from_tensor_slices(series)
                                                     # Window the data but only take those with the specified size
                                                     dataset = dataset.window(window_size + 1, shift=1, drop_remainder=True)
                                                     # Flatten the windows by putting its elements in a single batch
                                                     dataset = dataset.flat_map(lambda window: window.batch(window_size + 1))
                                                     # Create tuples with features and labels
                                                     dataset = dataset.map(lambda window: (window[:-1], window[-1]))
                                                     # Shuffle the windows
                                                     dataset = dataset.shuffle(shuffle_buffer)
                                                     # Create batches of windows
                                                     dataset = dataset.batch(batch_size)
                                                     # Optimize the dataset for training
                                                     dataset = dataset.cache().prefetch(1)
                                                     return dataset
```

As mentioned in the lectures, if your results don't look good, you can try tweaking the parameters here and see if the model will learn better.

```
In [7]: # Parameters
    window_size = 30
    batch_size = 32
    shuffle_buffer_size = 1000

# Generate the dataset windows
    train_set = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)
```

Build the Model

You've seen these layers before and here is how it looks like when combined.

```
# Print the model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 30, 64)	256
lstm (LSTM)	(None, 30, 64)	33,024
lstm_1 (LSTM)	(None, 64)	33,024
dense (Dense)	(None, 30)	1,950
dense_1 (Dense)	(None, 10)	310
dense_2 (Dense)	(None, 1)	11
lambda (Lambda)	(None, 1)	0

Total params: 68,575 (267.87 KB)

Trainable params: 68,575 (267.87 KB)

Non-trainable params: 0 (0.00 B)

Tune the Learning Rate

As usual, you will want to pick an optimal learning rate.

```
In [9]: # Get initial weights
    init_weights = model.get_weights()

In [10]: # Set the Learning rate scheduler
    lr_schedule = tf.keras.callbacks.LearningRateScheduler(
        lambda epoch: le-8 * 10**(epoch / 20))

# Initialize the optimizer
    optimizer = tf.keras.optimizers.SGD(momentum=0.9)

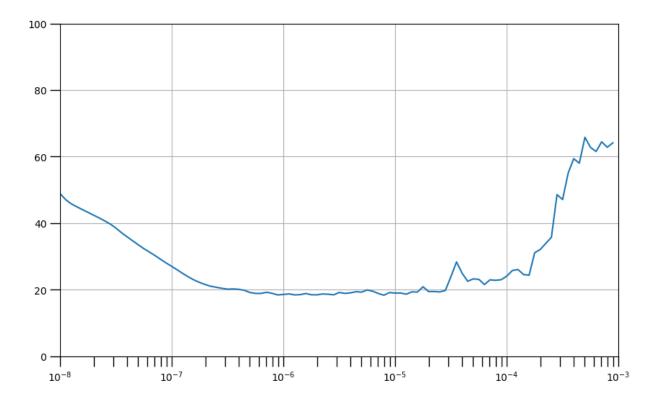
# Set the training parameters
    model.compile(loss=tf.keras.losses.Huber(), optimizer=optimizer)

# Train the model
    history = model.fit(train_set, epochs=100, callbacks=[lr_schedule])
```

Epoch 93/93	1/100	3s	5ms/sten	_	loss:	46.6301	_	learning_rate: 1.0000e-08
Epoch	2/100							
93/93 Epoch	3/100	0s	4ms/step	-	loss:	44.9357	-	learning_rate: 1.1220e-08
93/93 Enoch	4/100	0s	4ms/step	-	loss:	43.5745	-	learning_rate: 1.2589e-08
93/93		0s	4ms/step	-	loss:	42.6802	-	learning_rate: 1.4125e-08
Epoch 93/93	5/100	0s	4ms/step	_	loss:	41.8557	_	learning_rate: 1.5849e-08
Epoch 93/93	6/100	Ac	Ams/stan	_	1000	<i>1</i> 1 0677	_	learning rate: 1.7783e-08
Epoch	7/100							<u>-</u>
93/93 Epoch	8/100	0s	4ms/step	-	loss:	40.2823	-	learning_rate: 1.9953e-08
93/93 Enoch	9/100	0s	4ms/step	-	loss:	39.5021	-	learning_rate: 2.2387e-08
93/93		0s	4ms/step	-	loss:	38.7111	-	learning_rate: 2.5119e-08
Epoch 93/93	10/100	0s	4ms/step	_	loss:	37.8778	-	learning_rate: 2.8184e-08
Epoch 93/93	11/100	05	4ms/sten	_	loss:	36.9131	_	<pre>learning_rate: 3.1623e-08</pre>
Epoch	12/100							
	13/100							learning_rate: 3.5481e-08
93/93 Epoch	14/100	0s	4ms/step	-	loss:	34.6146	-	learning_rate: 3.9811e-08
	15/100	0s	4ms/step	-	loss:	33.5262	-	learning_rate: 4.4668e-08
93/93		0s	4ms/step	-	loss:	32.4430	-	learning_rate: 5.0119e-08
Epoch 93/93	16/100	0s	4ms/step	-	loss:	31.3676	-	learning_rate: 5.6234e-08
Epoch 93/93	17/100	0s	4ms/step	_	loss:	30.4817	_	learning_rate: 6.3096e-08
	18/100							learning_rate: 7.0795e-08
Epoch	19/100							
93/93 Epoch	20/100	05	4ms/step	-	loss:	28.5688	-	learning_rate: 7.9433e-08
93/93 Epoch	21/100	0s	4ms/step	-	loss:	27.4822	-	learning_rate: 8.9125e-08
93/93		0s	4ms/step	-	loss:	26.5376	-	learning_rate: 1.0000e-07
93/93		0s	4ms/step	-	loss:	25.5747	-	learning_rate: 1.1220e-07
Epoch 93/93	23/100	0s	4ms/step	-	loss:	24.6172	-	learning_rate: 1.2589e-07
Epoch 93/93	24/100	0s	4ms/step	_	loss:	23.6779	_	learning_rate: 1.4125e-07
	25/100							learning rate: 1.5849e-07
Epoch	26/100							<u>-</u>
93/93 Epoch	27/100	0s	4ms/step	-	loss:	22.1112	-	learning_rate: 1.7783e-07
93/93 Enoch	28/100	0s	4ms/step	-	loss:	21.5032	-	learning_rate: 1.9953e-07
93/93		0s	4ms/step	-	loss:	20.9637	-	learning_rate: 2.2387e-07
Еросп 93/93	29/100	0s	4ms/step	-	loss:	20.5714	-	learning_rate: 2.5119e-07
	30/100	0s	4ms/step	_	loss:	20.1957	_	learning_rate: 2.8184e-07
Epoch 93/93	31/100							learning_rate: 3.1623e-07
Epoch	32/100							
	33/100							learning_rate: 3.5481e-07
93/93 Epoch	34/100	0s	4ms/step	-	loss:	19.8676	-	learning_rate: 3.9811e-07
93/93		0s	4ms/step	-	loss:	19.4121	-	learning_rate: 4.4668e-07
93/93		0s	4ms/step	-	loss:	19.1286	-	learning_rate: 5.0119e-07
Epoch 93/93	36/100	0s	4ms/step	-	loss:	18.9326	-	learning_rate: 5.6234e-07
Epoch 93/93	37/100	0s	4ms/step	_	loss:	18.8238	_	learning_rate: 6.3096e-07
	38/100							learning_rate: 7.0795e-07
Epoch	39/100							
93/93		0s	4ms/step	-	loss:	18.8190	-	learning_rate: 7.9433e-07

Epoch	40/100								
93/93 Epoch	41/100	0s	4ms/step	-	loss:	18.3363	-	learning_rate:	8.9125e-07
93/93		0s	4ms/step	-	loss:	18.4176	-	<pre>learning_rate:</pre>	1.0000e-06
93/93	42/100	0s	4ms/step	-	loss:	18.7096	-	learning_rate:	1.1220e-06
Epoch 93/93	43/100	95	4ms/sten	_	loss:	18.3454	_	learning_rate:	1.2589e-06
Epoch	44/100								
93/93 Epoch	45/100	0s	4ms/step	-	loss:	18.331/	-	learning_rate:	1.4125e-06
93/93 Enoch	46/100	0s	4ms/step	-	loss:	18.9280	-	<pre>learning_rate:</pre>	1.5849e-06
93/93		0s	4ms/step	-	loss:	18.3650	-	<pre>learning_rate:</pre>	1.7783e-06
93/93	47/100	0s	4ms/step	-	loss:	18.3125	-	learning_rate:	1.9953e-06
Epoch 93/93	48/100	0s	4ms/step	_	loss:	18.4715	_	learning_rate:	2.2387e-06
	49/100							learning_rate:	
Epoch	50/100								
93/93 Epoch	51/100	0s	4ms/step	-	loss:	18.1129	-	learning_rate:	2.8184e-06
93/93 Enoch	52/100	0s	4ms/step	-	loss:	19.5236	-	<pre>learning_rate:</pre>	3.1623e-06
93/93		0s	4ms/step	-	loss:	18.7135	-	<pre>learning_rate:</pre>	3.5481e-06
93/93	53/100	0s	4ms/step	-	loss:	18.5830	-	learning_rate:	3.9811e-06
Epoch 93/93	54/100	0s	4ms/step	_	loss:	19.7399	_	learning_rate:	4.4668e-06
Epoch 93/93	55/100							learning_rate:	
Epoch	56/100								
93/93 Epoch	57/100							learning_rate:	
93/93 Epoch	58/100	0s	4ms/step	-	loss:	19.7691	-	learning_rate:	6.3096e-06
93/93 Epoch	59/100	0s	4ms/step	-	loss:	18.6703	-	learning_rate:	7.0795e-06
93/93		0s	4ms/step	-	loss:	18.4371	-	<pre>learning_rate:</pre>	7.9433e-06
93/93		0s	4ms/step	-	loss:	19.3801	-	learning_rate:	8.9125e-06
93/93		0s	4ms/step	-	loss:	18.8798	-	learning_rate:	1.0000e-05
Epoch 93/93	62/100	0s	4ms/step	-	loss:	19.4085	-	learning_rate:	1.1220e-05
Epoch 93/93	63/100	0s	4ms/step	_	loss:	18.5000	_	learning_rate:	1.2589e-05
Epoch 93/93	64/100							learning_rate:	
Epoch	65/100							learning rate:	
-	66/100							0_	
93/93 Epoch	67/100	0s	4ms/step	-	loss:	20.5147	-	learning_rate:	1.7783e-05
93/93 Epoch	68/100	0s	4ms/step	-	loss:	19.4624	-	learning_rate:	1.9953e-05
93/93 Enoch	69/100	0s	4ms/step	-	loss:	19.4803	-	<pre>learning_rate:</pre>	2.2387e-05
93/93		0s	4ms/step	-	loss:	20.0684	-	<pre>learning_rate:</pre>	2.5119e-05
93/93		0s	4ms/step	-	loss:	19.1796	-	learning_rate:	2.8184e-05
Epoch 93/93	71/100	0s	4ms/step	_	loss:	21.7473	_	learning_rate:	3.1623e-05
Epoch 93/93	72/100	0s	4ms/step	_	loss:	29.4175	_	<pre>learning_rate:</pre>	3.5481e-05
	73/100							learning_rate:	
Epoch	74/100								
-	75/100							learning_rate:	
93/93 Epoch	76/100	0s	4ms/step	-	loss:	21.2625	-	learning_rate:	5.0119e-05
93/93 Epoch	77/100	0s	4ms/step	-	loss:	24.1812	-	<pre>learning_rate:</pre>	5.6234e-05
93/93		0s	4ms/step	-	loss:	21.4995	-	<pre>learning_rate:</pre>	6.3096e-05
93/93		0s	4ms/step	-	loss:	21.3179	-	<pre>learning_rate:</pre>	7.0795e-05

```
Epoch 79/100
        93/93 •
                                   0s 4ms/step - loss: 22.8117 - learning_rate: 7.9433e-05
        Enoch 80/100
        93/93
                                   0s 4ms/step - loss: 22.3768 - learning_rate: 8.9125e-05
        Epoch 81/100
        93/93 •
                                   0s 4ms/step - loss: 22.3952 - learning_rate: 1.0000e-04
        Epoch 82/100
        93/93
                                   0s 4ms/step - loss: 25.6412 - learning_rate: 1.1220e-04
        Epoch 83/100
        93/93
                                  - 0s 4ms/step - loss: 24.4390 - learning_rate: 1.2589e-04
        Epoch 84/100
        93/93
                                  - 0s 4ms/step - loss: 23.3215 - learning_rate: 1.4125e-04
        Epoch 85/100
        93/93 •
                                  - 0s 4ms/step - loss: 24.1262 - learning_rate: 1.5849e-04
        Epoch 86/100
        93/93
                                  - 0s 4ms/step - loss: 30.2247 - learning_rate: 1.7783e-04
        Epoch 87/100
        93/93
                                  - 0s 4ms/step - loss: 33.5329 - learning_rate: 1.9953e-04
        Epoch 88/100
                                  - 0s 4ms/step - loss: 33.6691 - learning_rate: 2.2387e-04
        93/93
        Epoch 89/100
                                   0s 4ms/step - loss: 31.6961 - learning_rate: 2.5119e-04
        93/93
        Epoch 90/100
                                   0s 4ms/step - loss: 43.0452 - learning_rate: 2.8184e-04
        93/93
        Epoch 91/100
        93/93
                                  - 0s 4ms/step - loss: 47.1326 - learning_rate: 3.1623e-04
        Epoch 92/100
        93/93
                                  - 0s 4ms/step - loss: 51.0992 - learning_rate: 3.5481e-04
        Epoch 93/100
        93/93
                                  - 0s 4ms/step - loss: 58.3860 - learning_rate: 3.9811e-04
        Epoch 94/100
        93/93
                                  - 0s 4ms/step - loss: 56.6061 - learning_rate: 4.4668e-04
        Epoch 95/100
        93/93
                                  - 0s 4ms/step - loss: 61.0482 - learning rate: 5.0119e-04
        Epoch 96/100
        93/93 •
                                   0s 4ms/step - loss: 63.5604 - learning_rate: 5.6234e-04
        Epoch 97/100
                                  - 0s 4ms/step - loss: 60.2895 - learning_rate: 6.3096e-04
        93/93
        Epoch 98/100
        93/93
                                  Os 4ms/step - loss: 63.9995 - learning_rate: 7.0795e-04
        Epoch 99/100
        93/93
                                  • 0s 4ms/step - loss: 57.3688 - learning_rate: 7.9433e-04
        Epoch 100/100
        93/93
                                  - 0s 4ms/step - loss: 62.2880 - learning_rate: 8.9125e-04
In [11]: # Define the learning rate array
         lrs = 1e-8 * (10 ** (np.arange(100) / 20))
         # Set the figure size
         plt.figure(figsize=(10, 6))
         # Set the grid
         plt.grid(True)
         # Plot the loss in log scale
         \verb|plt.semilogx(lrs, history.history["loss"])| \\
         # Increase the tickmarks size
         plt.tick_params('both', length=10, width=1, which='both')
         # Set the plot boundaries
         plt.axis([1e-8, 1e-3, 0, 100])
Out[11]: (1e-08, 0.001, 0.0, 100.0)
```



Train the Model

Now you can proceed to reset and train the model. It is set for 100 epochs in the cell below but feel free to increase it if you want. Laurence got his results in the lectures after 500.

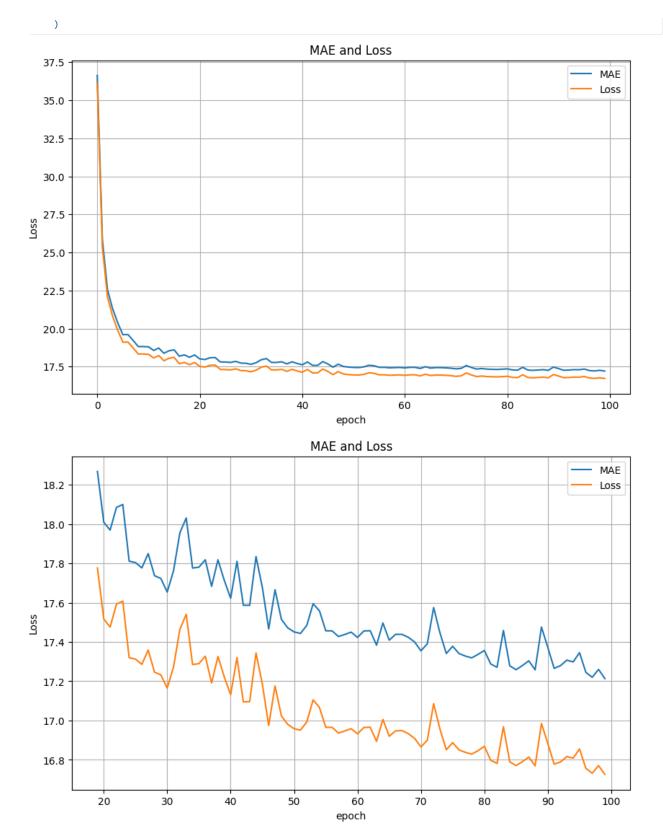
Epoch									
93/93 Epoch	2/100	2s	4ms/step	-	loss:	39.6136	-	mae:	40.1081
93/93		0s	4ms/step	-	loss:	25.9470	-	mae:	26.4412
Epoch 93/93		0s	4ms/step	_	loss:	21.9388	_	mae:	22.4321
Epoch	4/100								
93/93 Epoch		0s	4ms/step	-	loss:	20.5369	-	mae:	21.0293
93/93		0s	4ms/step	-	loss:	19.6171	-	mae:	20.1084
Epoch 93/93	6/100	0s	4ms/step	_	loss:	19.0719	_	mae:	19.5636
-	7/100	0-	1		1	10 0006			10 2010
93/93 Epoch	8/100	05	4ms/step	-	1055:	18.8906	-	mae:	19.3818
93/93 Enoch	9/100	0s	4ms/step	-	loss:	18.4057	-	mae:	18.8981
93/93		0s	4ms/step	-	loss:	18.3848	-	mae:	18.8761
Epoch 93/93	10/100	95	4ms/step	_	loss:	18.1387	_	mae:	18.6297
Epoch	11/100								
93/93 Epoch	12/100	0s	4ms/step	-	loss:	18.1928	-	mae:	18.6848
93/93		0s	4ms/step	-	loss:	17.9367	-	mae:	18.4281
93/93	13/100	0s	4ms/step	-	loss:	18.0193	-	mae:	18.5110
Epoch 93/93	14/100	۵s	4ms/step	_	loss	17 8187	_	mae.	18 3084
Epoch	15/100								
93/93 Epoch	16/100	0s	4ms/step	-	loss:	17.8371	-	mae:	18.3261
93/93		0s	4ms/step	-	loss:	18.0575	-	mae:	18.5479
93/93	17/100	0s	4ms/step	-	loss:	17.6693	-	mae:	18.1608
Epoch 93/93	18/100	0s	4ms/step	_	loss:	17.6179	_	mae:	18.1082
Epoch	19/100								
93/93 Epoch	20/100	05	4ms/step	-	1055:	17.5495	-	mae:	18.0394
93/93 Enoch	21/100	0s	4ms/step	-	loss:	17.7358	-	mae:	18.2272
93/93		0s	4ms/step	-	loss:	17.4467	-	mae:	17.9387
Epoch 93/93	22/100	0s	4ms/step	_	loss:	17.4295	_	mae:	17.9234
Epoch 93/93	23/100	۵c	4ms/step	_	1055	17 /1972		mae.	17 0707
	24/100								
93/93 Epoch	25/100	0s	4ms/step	-	loss:	17.5421	-	mae:	18.0328
93/93		0s	4ms/step	-	loss:	17.2417	-	mae:	17.7329
93/93	26/100	0s	4ms/step	-	loss:	17.2615	-	mae:	17.7527
Epoch 93/93	27/100	۵s	4ms/step	_	1000	17 231/	_	mae.	17 7220
Epoch	28/100								
93/93 Epoch	29/100	0s	4ms/step	-	loss:	17.2866	-	mae:	17.7775
93/93 Enach	30/100	0s	4ms/step	-	loss:	17.1964	-	mae:	17.6858
93/93		0s	4ms/step	-	loss:	17.1704	-	mae:	17.6596
Epoch 93/93	31/100	0s	4ms/step	_	loss:	17.1255	_	mae:	17.6134
Epoch	32/100								
93/93 Epoch	33/100	0 S	4ms/step	-	loss:	17.1940	-	mae:	17.6835
93/93 Enoch	34/100	0s	4ms/step	-	loss:	17.4106	-	mae:	17.9030
93/93		0s	4ms/step	-	loss:	17.4589	-	mae:	17.9484
Epoch 93/93	35/100	0s	4ms/step	_	loss:	17.2509	_	mae:	17.7417
Epoch	36/100								
93/93 Epoch	37/100	0 5	4ms/step	-	1022;	17.2322	-	mae:	1/./241
93/93 Epoch	38/100	0s	4ms/step	-	loss:	17.2800	-	mae:	17.7701
93/93		0s	4ms/step	-	loss:	17.1179	-	mae:	17.6078
Epoch 93/93	39/100	0s	4ms/step	-	loss:	17.2415	-	mae:	17.7345

Epoch	40/100								
93/93		0s	4ms/step	-	loss:	17.1465	-	mae:	17.6383
	41/100	00	1mc/c+on		10001	17 0101			17 5000
93/93 Epoch	42/100	05	4ms/step	-	1055;	17.0101	-	mae:	17.5009
93/93		0s	4ms/step	-	loss:	17.3520	-	mae:	17.8417
	43/100	_			,	.=			.=
93/93 Enoch	44/100	0s	4ms/step	-	loss:	17.0269	-	mae:	17.5194
93/93		0s	4ms/step	_	loss:	17.0558	_	mae:	17.5469
Epoch	45/100								
93/93		0s	4ms/step	-	loss:	17.2432	-	mae:	17.7361
93/93	46/100	0s	4ms/step	_	loss:	17.1332	_	mae:	17,6261
	47/100		,						
93/93		0s	4ms/step	-	loss:	16.9260	-	mae:	17.4170
93/93	48/100	05	4ms/step	_	loss:	17.1361	_	mae:	17.6267
	49/100		э, эсср		1055.	17,12301			1,,020,
93/93		0s	4ms/step	-	loss:	16.9581	-	mae:	17.4505
Epoch 93/93	50/100	۵c	4ms/step	_	1000	16 8992	_	mae.	17 3913
	51/100	03	4 11137 3 сср		1033.	10.0552		mac.	17.3313
		0s	4ms/step	-	loss:	16.8902	-	mae:	17.3835
Epoch 93/93	52/100	۵c	4ms/step	_	1000	16 9077	_	mae.	17 3999
	53/100	03	-1113/3cep	-	1033.	10.90//	-	muc.	11.0000
93/93		0s	4ms/step	-	loss:	16.8711	-	mae:	17.3607
Epoch 93/93	54/100	Q.c	4ms/step		1000	17 0090		m20:	17 E001
	55/100	05	41115/5CEP	-	1055.	17.0500	-	mae.	17.3001
93/93		0s	4ms/step	-	loss:	17.1037	-	mae:	17.5965
	56/100	00	1mc/c+on		10001	16 0014			17 4005
93/93 Epoch	57/100	05	4ms/step	-	1055:	10.9914	-	mae:	17.4625
93/93		0s	4ms/step	-	loss:	16.8659	-	mae:	17.3571
	58/100	00	1mc/c+on		10001	16 0074			17 2002
93/93 Epoch	59/100	05	4ms/step	-	1055:	16.9074	-	mae:	17.3993
93/93		0s	4ms/step	-	loss:	16.8539	-	mae:	17.3454
	60/100	0-	4/		1	16 0727			17 4645
93/93 Epoch	61/100	05	4ms/step	-	1055:	16.9/2/	-	mae:	17.4645
93/93		0s	4ms/step	-	loss:	16.8607	-	mae:	17.3511
	62/100	00	1mc/c+on		10001	17 0001			17 5017
93/93 Epoch	63/100	05	4ms/step	-	1055:	17.0091	-	mae:	17.5017
93/93		0s	4ms/step	-	loss:	16.9475	-	mae:	17.4388
	64/100	0-	4/		1	16 0020			17 2025
93/93 Epoch	65/100	05	4ms/step	-	1055:	10.0939	-	mae:	17.3033
93/93		0s	4ms/step	-	loss:	16.9111	-	mae:	17.4029
	66/100	0-	4/		1	16 0026			17 2015
93/93 Epoch	67/100	05	4ms/step	-	1055:	16.9026	-	mae:	17.3915
93/93		0s	4ms/step	-	loss:	16.8482	-	mae:	17.3393
	68/100	0-	4/		1	16 0241			17 4240
93/93 Epoch	69/100	05	4ms/step	-	1055:	10.9341	-	mae:	17.4240
		0s	4ms/step	-	loss:	16.9285	-	mae:	17.4198
Epoch 93/93	70/100	00	1mc/c+on		10001	16 0020			17 2044
	71/100	05	4ms/step	-	1055:	16.8928	-	mae:	17.3844
		0s	4ms/step	-	loss:	16.7951	-	mae:	17.2840
	72/100	0-	4/		1	16 0560			17 2475
93/93 Epoch	73/100	05	4ms/step	-	1055:	16.8568	-	mae:	17.34/5
		0s	4ms/step	-	loss:	16.9939	-	mae:	17.4838
•	74/100	0-	1mc / c+ c=		locs	16 0500		m 2 0 :	17 2400
93/93 Epoch	75/100	0 5	4ms/step	-	1022:	10.8299	-	mae:	17.3498
93/93		0s	4ms/step	-	loss:	16.7715	-	mae:	17.2613
	76/100	0-	1mc / = 1 -		16	16 0052			17 2070
93/93 Epoch	77/100	ØS	4ms/step	-	TO22:	10.8023	-	mae:	17.2970
93/93		0s	4ms/step	-	loss:	16.7764	-	mae:	17.2664
	78/100	00	1mc/c+00	_	1055	16 7//1	_	mac.	17 2242
93/93		05	4ms/step	-	1022;	10./441	-	mae:	11.2343

```
Epoch 79/100
93/93
                          0s 4ms/step - loss: 16.7577 - mae: 17.2491
Enoch 80/100
93/93
                          0s 4ms/step - loss: 16.7868 - mae: 17.2770
Epoch 81/100
93/93
                          0s 4ms/step - loss: 16.9254 - mae: 17.4148
Epoch 82/100
93/93
                          0s 4ms/step - loss: 16.7296 - mae: 17.2204
Epoch 83/100
93/93
                          0s 4ms/step - loss: 16.6958 - mae: 17.1842
Epoch 84/100
93/93
                          0s 4ms/step - loss: 16.8964 - mae: 17.3877
Epoch 85/100
93/93
                          0s 4ms/step - loss: 16.7164 - mae: 17.2038
Epoch 86/100
93/93
                          0s 4ms/step - loss: 16.7003 - mae: 17.1883
Epoch 87/100
                          0s 4ms/step - loss: 16.6903 - mae: 17.1794
93/93
Epoch 88/100
                          0s 4ms/step - loss: 16.7908 - mae: 17.2813
93/93
Epoch 89/100
                          0s 4ms/step - loss: 16.6933 - mae: 17.1829
93/93
Epoch 90/100
93/93
                          0s 4ms/step - loss: 16.8816 - mae: 17.3725
Epoch 91/100
93/93
                          0s 4ms/step - loss: 16.8006 - mae: 17.2887
Epoch 92/100
93/93
                          0s 4ms/step - loss: 16.7060 - mae: 17.1929
Epoch 93/100
93/93
                          0s 4ms/step - loss: 16.7256 - mae: 17.2167
Epoch 94/100
93/93
                          0s 4ms/step - loss: 16.7363 - mae: 17.2279
Epoch 95/100
93/93
                          0s 4ms/step - loss: 16.7158 - mae: 17.2047
Epoch 96/100
93/93
                          0s 4ms/step - loss: 16.8175 - mae: 17.3087
Epoch 97/100
                          0s 4ms/step - loss: 16.6731 - mae: 17.1622
93/93
Epoch 98/100
93/93
                          0s 4ms/step - loss: 16.6614 - mae: 17.1494
Epoch 99/100
93/93
                          0s 4ms/step - loss: 16.6851 - mae: 17.1756
Epoch 100/100
93/93
                          0s 4ms/step - loss: 16.6527 - mae: 17.1410
```

You can visualize the training and see if the loss and MAE are still trending down.

```
In [15]: # Get mae and loss from history log
         mae=history.history['mae']
         loss=history.history['loss']
         # Get number of epochs
         epochs=range(len(loss))
         # Plot mae and loss
         plot_series(
             x=epochs,
             y=(mae, loss),
             title='MAE and Loss',
             xlabel='epoch',
             ylabel='Loss',
             legend=['MAE', 'Loss']
         # Only plot the last 80% of the epochs
         zoom_split = int(epochs[-1] * 0.2)
         epochs_zoom = epochs[zoom_split:]
         mae zoom = mae[zoom split:]
         loss_zoom = loss[zoom_split:]
         # Plot zoomed mae and loss
         plot_series(
             x=epochs_zoom,
             y=(mae_zoom, loss_zoom),
             title='MAE and Loss',
             xlabel='epoch',
             ylabel='Loss',
             legend=['MAE', 'Loss']
```



Model Prediction

As before, you can get the predictions for the validation set time range and compute the metrics.

```
In [16]: def model_forecast(model, series, window_size, batch_size):
    """Uses an input model to generate predictions on data windows

Args:
    model (TF Keras Model) - model that accepts data windows
```

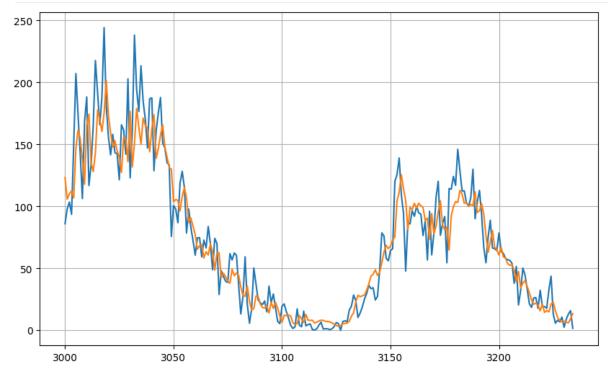
```
series (array of float) - contains the values of the time series
 window_size (int) - the number of time steps to include in the window
 batch\_size (int) - the batch size
Returns:
 forecast (numpy array) - array containing predictions
# Add an axis for the feature dimension of RNN Layers
series = tf.expand_dims(series, axis=-1)
# Generate a TF Dataset from the series values
dataset = tf.data.Dataset.from tensor slices(series)
# Window the data but only take those with the specified size
dataset = dataset.window(window_size, shift=1, drop_remainder=True)
# Flatten the windows by putting its elements in a single batch
dataset = dataset.flat_map(lambda w: w.batch(window_size))
# Create batches of windows
dataset = dataset.batch(batch_size).prefetch(1)
# Get predictions on the entire dataset
forecast = model.predict(dataset, verbose=0)
return forecast
```

```
In [17]: # Reduce the original series
    forecast_series = series[split_time-window_size:-1]

# Use helper function to generate predictions
    forecast = model_forecast(model, forecast_series, window_size, batch_size)

# Drop single dimensional axis
    results = forecast.squeeze()

# Plot the results
    plot_series(time_valid, (x_valid, results))
```



```
In [18]: # Compute the MAE
    print(tf.keras.metrics.mae(x_valid, results).numpy())
```

14.527155

This concludes the final practice lab for this course! You implemented a deep and complex architecture composed of CNNs, RNNs, and
DNNs. You'll be using the skills you developed throughout this course to complete the final assignment. Keep it up!