

# Ungraded Lab: Predicting Sunspots with Neural Networks

In the remaining labs for this week, you will move away from synthetic time series and start building models for real world data. In particular, you will train on the [Sunspots](#) dataset: a monthly record of sunspot numbers from January 1749 to July 2018. You will first build a deep neural network here composed of dense layers. This will act as your baseline so you can compare it to the next lab where you will use a more complex architecture.

Let's begin!

## Imports

You will use the same imports as before with the addition of the `csv` module. You will need this to parse the CSV file containing the dataset.

```
In [1]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import csv
```

## Utilities

You will only have the `plot_series()` dataset here because you no longer need the synthetic data generation functions.

```
In [2]: def plot_series(x, y, format="-", start=0, end=None,
                        title=None, xlabel=None, ylabel=None, legend=None ):
    """
    Visualizes time series data

    Args:
        x (array of int) - contains values for the x-axis
        y (array of int or tuple of arrays) - contains the values for the y-axis
        format (string) - line style when plotting the graph
        label (string) - tag for the line
        start (int) - first time step to plot
        end (int) - last time step to plot
        title (string) - title of the plot
        xlabel (string) - label for the x-axis
        ylabel (string) - label for the y-axis
        legend (list of strings) - legend for the plot
    """

    # Setup dimensions of the graph figure
    plt.figure(figsize=(10, 6))

    # Check if there are more than two series to plot
    if type(y) is tuple:

        # Loop over the y elements
        for y_curr in y:

            # Plot the x and current y values
            plt.plot(x[start:end], y_curr[start:end], format)

    else:
        # Plot the x and y values
        plt.plot(x[start:end], y[start:end], format)

    # Label the x-axis
    plt.xlabel(xlabel)

    # Label the y-axis
    plt.ylabel(ylabel)

    # Set the Legend
    if legend:
        plt.legend(legend)

    # Set the title
    plt.title(title)

    # Overlay a grid on the graph
    plt.grid(True)

    # Draw the graph on screen
    plt.show()
```

## Download and Preview the Dataset

You can now download the dataset and inspect the contents. The link in class is from Laurence's repo but we also hosted it in the link below.

```
In [3]: # Download the dataset
!wget -nc https://storage.googleapis.com/tensorflow-1-public/course4/Sunspots.csv
```

File 'Sunspots.csv' already there; not retrieving.

Running the cell below, you'll see that there are only three columns in the dataset:

1. untitled column containing the month number
2. Date which has the format YYYY-MM-DD
3. Mean Total Sunspot Number

```
In [4]: # Preview the dataset
!head Sunspots.csv
```

```
,Date,Monthly Mean Total Sunspot Number
0,1749-01-31,96.7
1,1749-02-28,104.3
2,1749-03-31,116.7
3,1749-04-30,92.8
4,1749-05-31,141.7
5,1749-06-30,139.2
6,1749-07-31,158.0
7,1749-08-31,110.5
8,1749-09-30,126.5
```

For this lab and the next, you will only need the month number and the mean total sunspot number. You will load those into memory and convert it to arrays that represents a time series.

```
In [5]: # Initialize Lists
time_step = []
sunspots = []

# Open CSV file
with open('./Sunspots.csv') as csvfile:

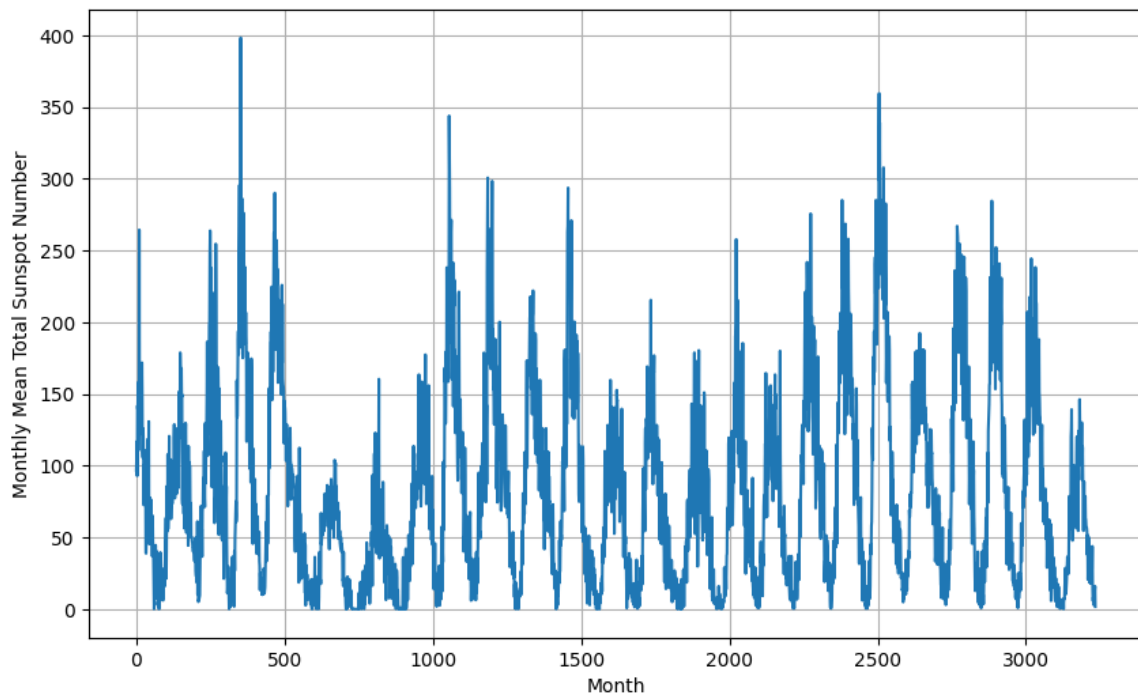
    # Initialize reader
    reader = csv.reader(csvfile, delimiter=',')

    # Skip the first line
    next(reader)

    # Append row and sunspot number to Lists
    for row in reader:
        time_step.append(int(row[0]))
        sunspots.append(float(row[2]))

# Convert Lists to numpy arrays
time = np.array(time_step)
series = np.array(sunspots)

# Preview the data
plot_series(time, series, xlabel='Month', ylabel='Monthly Mean Total Sunspot Number')
```



## Split the Dataset

Next, you will split the dataset into training and validation sets. There are 3235 points in the dataset and you will use the first 3000 for training.

```
In [6]: # Define the split time
split_time = 3000

# Get the train set
time_train = time[:split_time]
x_train = series[:split_time]

# Get the validation set
time_valid = time[split_time:]
x_valid = series[split_time:]
```

## Prepare Features and Labels

You can then prepare the dataset windows as before. The window size is set to 30 points (equal to 2.5 years) but feel free to change later on if you want to experiment.

```
In [7]: def windowed_dataset(series, window_size, batch_size, shuffle_buffer):
        """Generates dataset windows

        Args:
            series (array of float) - contains the values of the time series
            window_size (int) - the number of time steps to include in the feature
            batch_size (int) - the batch size
            shuffle_buffer(int) - buffer size to use for the shuffle method

        Returns:
            dataset (TF Dataset) - TF Dataset containing time windows
        """

        # Generate a TF Dataset from the series values
        dataset = tf.data.Dataset.from_tensor_slices(series)

        # Window the data but only take those with the specified size
        dataset = dataset.window(window_size + 1, shift=1, drop_remainder=True)

        # Flatten the windows by putting its elements in a single batch
        dataset = dataset.flat_map(lambda window: window.batch(window_size + 1))

        # Create tuples with features and labels
        dataset = dataset.map(lambda window: (window[:-1], window[-1]))

        # Shuffle the windows
        dataset = dataset.shuffle(shuffle_buffer)
```

```

# Create batches of windows
dataset = dataset.batch(batch_size)

# Optimize the dataset for training
dataset = dataset.cache().prefetch(1)

return dataset

```

```

In [8]: # Parameters
window_size = 30
batch_size = 32
shuffle_buffer_size = 1000

# Generate the dataset windows
train_set = windowed_dataset(x_train, window_size, batch_size, shuffle_buffer_size)

```

## Build the Model

The model will be 3-layer dense network as shown below.

```

In [9]: # Build the model
model = tf.keras.models.Sequential([
    tf.keras.Input(shape=(window_size,)),
    tf.keras.layers.Dense(30, activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1)
])

# Print the model summary
model.summary()

```

Model: "sequential"

| Layer (type)    | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense)   | (None, 30)   | 930     |
| dense_1 (Dense) | (None, 10)   | 310     |
| dense_2 (Dense) | (None, 1)    | 11      |

Total params: 1,251 (4.89 KB)

Trainable params: 1,251 (4.89 KB)

Non-trainable params: 0 (0.00 B)

## Tune the Learning Rate

You can pick a learning rate by running the same learning rate scheduler code from previous labs.

```

In [10]: # Set the Learning rate scheduler
lr_schedule = tf.keras.callbacks.LearningRateScheduler(
    lambda epoch: 1e-8 * 10**(epoch / 20))

# Initialize the optimizer
optimizer = tf.keras.optimizers.SGD(momentum=0.9)

# Set the training parameters
model.compile(loss=tf.keras.losses.Huber(), optimizer=optimizer)

# Train the model
history = model.fit(train_set, epochs=100, callbacks=[lr_schedule])

```

Epoch 1/100

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR










































I0000 00:00:1745446178.871988 30792 service.cc:145] XLA service 0x7c110000a230 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1745446178.872825 30792 service.cc:153] StreamExecutor device (0): NVIDIA A10G, Compute Capability 8.6

51/Unknown 1s 3ms/step - loss: 53.3019

I0000 00:00:1745446179.579299 30792 device\_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

93/93 ————— 2s 6ms/step - loss: 55.1225 - learning\_rate: 1.0000e-08  
Epoch 2/100  
93/93 ————— 0s 751us/step - loss: 54.6505 - learning\_rate: 1.1220e-08  
Epoch 3/100  
93/93 ————— 0s 768us/step - loss: 54.1112 - learning\_rate: 1.2589e-08  
Epoch 4/100  
93/93 ————— 0s 735us/step - loss: 53.5123 - learning\_rate: 1.4125e-08  
Epoch 5/100  
93/93 ————— 0s 724us/step - loss: 52.8496 - learning\_rate: 1.5849e-08  
Epoch 6/100  
93/93 ————— 0s 738us/step - loss: 52.1164 - learning\_rate: 1.7783e-08  
Epoch 7/100  
93/93 ————— 0s 734us/step - loss: 51.3071 - learning\_rate: 1.9953e-08  
Epoch 8/100  
93/93 ————— 0s 726us/step - loss: 50.4076 - learning\_rate: 2.2387e-08  
Epoch 9/100  
93/93 ————— 0s 771us/step - loss: 49.4352 - learning\_rate: 2.5119e-08  
Epoch 10/100  
93/93 ————— 0s 762us/step - loss: 48.3740 - learning\_rate: 2.8184e-08  
Epoch 11/100  
93/93 ————— 0s 748us/step - loss: 47.2309 - learning\_rate: 3.1623e-08  
Epoch 12/100  
93/93 ————— 0s 752us/step - loss: 46.0027 - learning\_rate: 3.5481e-08  
Epoch 13/100  
93/93 ————— 0s 740us/step - loss: 44.6798 - learning\_rate: 3.9811e-08  
Epoch 14/100  
93/93 ————— 0s 727us/step - loss: 43.2797 - learning\_rate: 4.4668e-08  
Epoch 15/100  
93/93 ————— 0s 736us/step - loss: 41.8520 - learning\_rate: 5.0119e-08  
Epoch 16/100  
93/93 ————— 0s 738us/step - loss: 40.4045 - learning\_rate: 5.6234e-08  
Epoch 17/100  
93/93 ————— 0s 721us/step - loss: 38.9533 - learning\_rate: 6.3096e-08  
Epoch 18/100  
93/93 ————— 0s 722us/step - loss: 37.5629 - learning\_rate: 7.0795e-08  
Epoch 19/100  
93/93 ————— 0s 750us/step - loss: 36.2016 - learning\_rate: 7.9433e-08  
Epoch 20/100  
93/93 ————— 0s 751us/step - loss: 34.9788 - learning\_rate: 8.9125e-08  
Epoch 21/100  
93/93 ————— 0s 745us/step - loss: 33.9053 - learning\_rate: 1.0000e-07  
Epoch 22/100  
93/93 ————— 0s 741us/step - loss: 32.9742 - learning\_rate: 1.1220e-07  
Epoch 23/100  
93/93 ————— 0s 743us/step - loss: 32.2089 - learning\_rate: 1.2589e-07  
Epoch 24/100  
93/93 ————— 0s 735us/step - loss: 31.5788 - learning\_rate: 1.4125e-07  
Epoch 25/100  
93/93 ————— 0s 738us/step - loss: 31.0358 - learning\_rate: 1.5849e-07  
Epoch 26/100  
93/93 ————— 0s 731us/step - loss: 30.5646 - learning\_rate: 1.7783e-07  
Epoch 27/100  
93/93 ————— 0s 729us/step - loss: 30.1131 - learning\_rate: 1.9953e-07  
Epoch 28/100  
93/93 ————— 0s 730us/step - loss: 29.6624 - learning\_rate: 2.2387e-07  
Epoch 29/100  
93/93 ————— 0s 722us/step - loss: 29.1867 - learning\_rate: 2.5119e-07  
Epoch 30/100  
93/93 ————— 0s 735us/step - loss: 28.6713 - learning\_rate: 2.8184e-07  
Epoch 31/100  
93/93 ————— 0s 737us/step - loss: 28.1073 - learning\_rate: 3.1623e-07  
Epoch 32/100  
93/93 ————— 0s 729us/step - loss: 27.4877 - learning\_rate: 3.5481e-07  
Epoch 33/100  
93/93 ————— 0s 738us/step - loss: 26.8408 - learning\_rate: 3.9811e-07  
Epoch 34/100  
93/93 ————— 0s 735us/step - loss: 26.1755 - learning\_rate: 4.4668e-07  
Epoch 35/100  
93/93 ————— 0s 752us/step - loss: 25.4903 - learning\_rate: 5.0119e-07  
Epoch 36/100  
93/93 ————— 0s 720us/step - loss: 24.7965 - learning\_rate: 5.6234e-07  
Epoch 37/100  
93/93 ————— 0s 734us/step - loss: 24.1260 - learning\_rate: 6.3096e-07  
Epoch 38/100  
93/93 ————— 0s 737us/step - loss: 23.4919 - learning\_rate: 7.0795e-07  
Epoch 39/100  
93/93 ————— 0s 729us/step - loss: 22.9322 - learning\_rate: 7.9433e-07  
Epoch 40/100  
93/93 ————— 0s 735us/step - loss: 22.4234 - learning\_rate: 8.9125e-07  
Epoch 41/100  
93/93 ————— 0s 713us/step - loss: 21.9676 - learning\_rate: 1.0000e-06  
Epoch 42/100  
93/93 ————— 0s 719us/step - loss: 21.5839 - learning\_rate: 1.1220e-06

Epoch 43/100  
93/93  0s 747us/step - loss: 21.2859 - learning\_rate: 1.2589e-06  
Epoch 44/100  
93/93  0s 736us/step - loss: 21.0300 - learning\_rate: 1.4125e-06  
Epoch 45/100  
93/93  0s 725us/step - loss: 20.8061 - learning\_rate: 1.5849e-06  
Epoch 46/100  
93/93  0s 719us/step - loss: 20.5994 - learning\_rate: 1.7783e-06  
Epoch 47/100  
93/93  0s 738us/step - loss: 20.4076 - learning\_rate: 1.9953e-06  
Epoch 48/100  
93/93  0s 728us/step - loss: 20.2202 - learning\_rate: 2.2387e-06  
Epoch 49/100  
93/93  0s 739us/step - loss: 20.0335 - learning\_rate: 2.5119e-06  
Epoch 50/100  
93/93  0s 730us/step - loss: 19.8553 - learning\_rate: 2.8184e-06  
Epoch 51/100  
93/93  0s 744us/step - loss: 19.6841 - learning\_rate: 3.1623e-06  
Epoch 52/100  
93/93  0s 744us/step - loss: 19.5060 - learning\_rate: 3.5481e-06  
Epoch 53/100  
93/93  0s 733us/step - loss: 19.3211 - learning\_rate: 3.9811e-06  
Epoch 54/100  
93/93  0s 750us/step - loss: 19.1599 - learning\_rate: 4.4668e-06  
Epoch 55/100  
93/93  0s 765us/step - loss: 18.9860 - learning\_rate: 5.0119e-06  
Epoch 56/100  
93/93  0s 761us/step - loss: 18.8047 - learning\_rate: 5.6234e-06  
Epoch 57/100  
93/93  0s 749us/step - loss: 18.6450 - learning\_rate: 6.3096e-06  
Epoch 58/100  
93/93  0s 774us/step - loss: 18.4664 - learning\_rate: 7.0795e-06  
Epoch 59/100  
93/93  0s 755us/step - loss: 18.3190 - learning\_rate: 7.9433e-06  
Epoch 60/100  
93/93  0s 772us/step - loss: 18.1975 - learning\_rate: 8.9125e-06  
Epoch 61/100  
93/93  0s 758us/step - loss: 18.0657 - learning\_rate: 1.0000e-05  
Epoch 62/100  
93/93  0s 729us/step - loss: 17.9759 - learning\_rate: 1.1220e-05  
Epoch 63/100  
93/93  0s 735us/step - loss: 17.8951 - learning\_rate: 1.2589e-05  
Epoch 64/100  
93/93  0s 724us/step - loss: 17.8090 - learning\_rate: 1.4125e-05  
Epoch 65/100  
93/93  0s 733us/step - loss: 17.7228 - learning\_rate: 1.5849e-05  
Epoch 66/100  
93/93  0s 727us/step - loss: 17.6416 - learning\_rate: 1.7783e-05  
Epoch 67/100  
93/93  0s 723us/step - loss: 17.5794 - learning\_rate: 1.9953e-05  
Epoch 68/100  
93/93  0s 703us/step - loss: 17.5761 - learning\_rate: 2.2387e-05  
Epoch 69/100  
93/93  0s 753us/step - loss: 17.4956 - learning\_rate: 2.5119e-05  
Epoch 70/100  
93/93  0s 760us/step - loss: 17.4184 - learning\_rate: 2.8184e-05  
Epoch 71/100  
93/93  0s 761us/step - loss: 17.4127 - learning\_rate: 3.1623e-05  
Epoch 72/100  
93/93  0s 748us/step - loss: 17.4420 - learning\_rate: 3.5481e-05  
Epoch 73/100  
93/93  0s 754us/step - loss: 17.3879 - learning\_rate: 3.9811e-05  
Epoch 74/100  
93/93  0s 743us/step - loss: 17.4239 - learning\_rate: 4.4668e-05  
Epoch 75/100  
93/93  0s 752us/step - loss: 17.4890 - learning\_rate: 5.0119e-05  
Epoch 76/100  
93/93  0s 755us/step - loss: 17.5174 - learning\_rate: 5.6234e-05  
Epoch 77/100  
93/93  0s 743us/step - loss: 17.9423 - learning\_rate: 6.3096e-05  
Epoch 78/100  
93/93  0s 739us/step - loss: 17.7957 - learning\_rate: 7.0795e-05  
Epoch 79/100  
93/93  0s 733us/step - loss: 17.9843 - learning\_rate: 7.9433e-05  
Epoch 80/100  
93/93  0s 734us/step - loss: 17.8700 - learning\_rate: 8.9125e-05  
Epoch 81/100  
93/93  0s 731us/step - loss: 17.4891 - learning\_rate: 1.0000e-04  
Epoch 82/100  
93/93  0s 734us/step - loss: 17.5478 - learning\_rate: 1.1220e-04  
Epoch 83/100  
93/93  0s 713us/step - loss: 17.5019 - learning\_rate: 1.2589e-04  
Epoch 84/100

```

93/93 ————— 0s 740us/step - loss: 17.6458 - learning_rate: 1.4125e-04
Epoch 85/100
93/93 ————— 0s 718us/step - loss: 17.5317 - learning_rate: 1.5849e-04
Epoch 86/100
93/93 ————— 0s 734us/step - loss: 17.5086 - learning_rate: 1.7783e-04
Epoch 87/100
93/93 ————— 0s 736us/step - loss: 17.7165 - learning_rate: 1.9953e-04
Epoch 88/100
93/93 ————— 0s 744us/step - loss: 17.7900 - learning_rate: 2.2387e-04
Epoch 89/100
93/93 ————— 0s 739us/step - loss: 18.2583 - learning_rate: 2.5119e-04
Epoch 90/100
93/93 ————— 0s 735us/step - loss: 18.6401 - learning_rate: 2.8184e-04
Epoch 91/100
93/93 ————— 0s 714us/step - loss: 21.8398 - learning_rate: 3.1623e-04
Epoch 92/100
93/93 ————— 0s 718us/step - loss: 18.1356 - learning_rate: 3.5481e-04
Epoch 93/100
93/93 ————— 0s 713us/step - loss: 18.4450 - learning_rate: 3.9811e-04
Epoch 94/100
93/93 ————— 0s 753us/step - loss: 20.2958 - learning_rate: 4.4668e-04
Epoch 95/100
93/93 ————— 0s 764us/step - loss: 22.2458 - learning_rate: 5.0119e-04
Epoch 96/100
93/93 ————— 0s 851us/step - loss: 21.4302 - learning_rate: 5.6234e-04
Epoch 97/100
93/93 ————— 0s 780us/step - loss: 21.7769 - learning_rate: 6.3096e-04
Epoch 98/100
93/93 ————— 0s 737us/step - loss: 20.0682 - learning_rate: 7.0795e-04
Epoch 99/100
93/93 ————— 0s 767us/step - loss: 23.7053 - learning_rate: 7.9433e-04
Epoch 100/100
93/93 ————— 0s 766us/step - loss: 24.6913 - learning_rate: 8.9125e-04

```

```

In [11]: # Define the Learning rate array
lrs = 1e-8 * (10 ** (np.arange(100) / 20))

# Set the figure size
plt.figure(figsize=(10, 6))

# Set the grid
plt.grid(True)

# Plot the loss in log scale
plt.semilogx(lrs, history.history["loss"])

# Increase the tickmarks size
plt.tick_params('both', length=10, width=1, which='both')

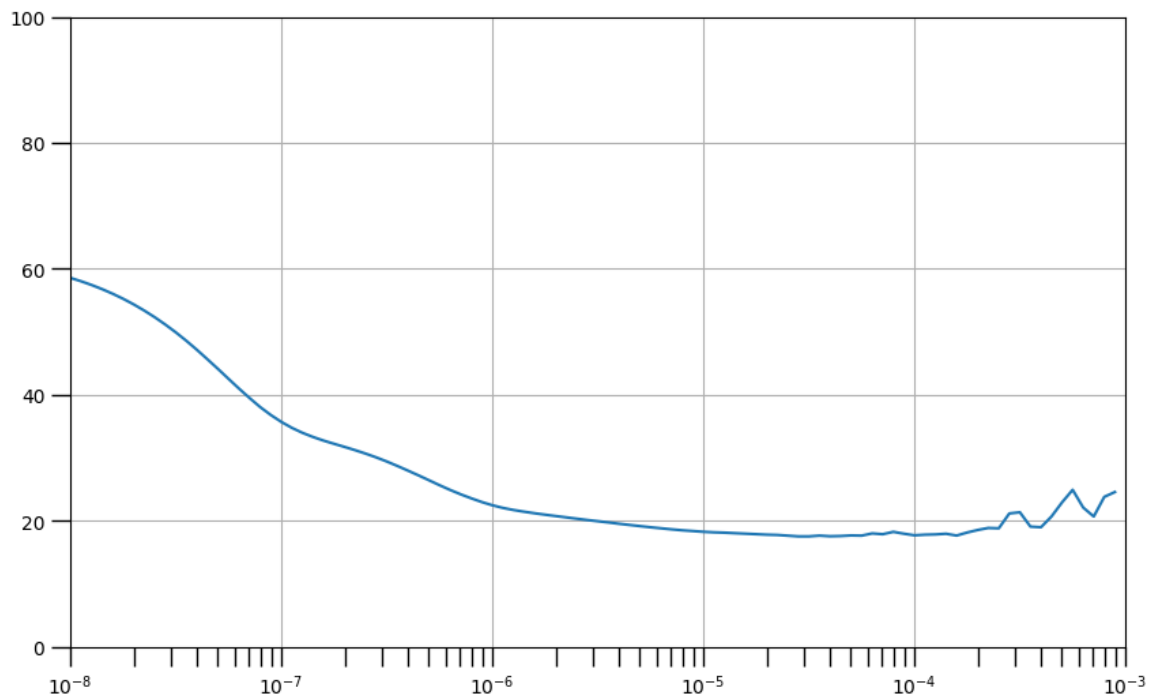
# Set the plot boundaries
plt.axis([1e-8, 1e-3, 0, 100])

```

```

Out[11]: (1e-08, 0.001, 0.0, 100.0)

```



## Train the Model

Once you've picked a learning rate, you can rebuild the model and start training.

```
In [12]: # Reset states generated by Keras
tf.keras.backend.clear_session()
```

```
# Build the Model
model = tf.keras.models.Sequential([
    tf.keras.Input(shape=(window_size,)),
    tf.keras.layers.Dense(30, activation="relu"),
    tf.keras.layers.Dense(10, activation="relu"),
    tf.keras.layers.Dense(1)
])
```

```
In [13]: # Set the Learning rate
learning_rate = 2e-5

# Set the optimizer
optimizer = tf.keras.optimizers.SGD(learning_rate=learning_rate, momentum=0.9)

# Set the training parameters
model.compile(loss=tf.keras.losses.Huber(),
              optimizer=optimizer,
              metrics=["mae"])

# Train the model
history = model.fit(train_set, epochs=100)
```



Epoch 1/100  
93/93 — 1s 4ms/step - loss: 27.6002 - mae: 28.0969  
Epoch 2/100  
93/93 — 0s 837us/step - loss: 19.9820 - mae: 20.4762  
Epoch 3/100  
93/93 — 0s 845us/step - loss: 18.7744 - mae: 19.2677  
Epoch 4/100  
93/93 — 0s 839us/step - loss: 18.2252 - mae: 18.7186  
Epoch 5/100  
93/93 — 0s 828us/step - loss: 17.9134 - mae: 18.4082  
Epoch 6/100  
93/93 — 0s 807us/step - loss: 17.6354 - mae: 18.1285  
Epoch 7/100  
93/93 — 0s 818us/step - loss: 17.4432 - mae: 17.9371  
Epoch 8/100  
93/93 — 0s 830us/step - loss: 17.3308 - mae: 17.8219  
Epoch 9/100  
93/93 — 0s 827us/step - loss: 17.2602 - mae: 17.7529  
Epoch 10/100  
93/93 — 0s 832us/step - loss: 17.1775 - mae: 17.6679  
Epoch 11/100  
93/93 — 0s 846us/step - loss: 17.0769 - mae: 17.5682  
Epoch 12/100  
93/93 — 0s 838us/step - loss: 17.0526 - mae: 17.5430  
Epoch 13/100  
93/93 — 0s 828us/step - loss: 16.9942 - mae: 17.4847  
Epoch 14/100  
93/93 — 0s 840us/step - loss: 16.9497 - mae: 17.4401  
Epoch 15/100  
93/93 — 0s 834us/step - loss: 16.9100 - mae: 17.4007  
Epoch 16/100  
93/93 — 0s 840us/step - loss: 16.8976 - mae: 17.3881  
Epoch 17/100  
93/93 — 0s 836us/step - loss: 16.8723 - mae: 17.3622  
Epoch 18/100  
93/93 — 0s 792us/step - loss: 16.8476 - mae: 17.3384  
Epoch 19/100  
93/93 — 0s 806us/step - loss: 16.8172 - mae: 17.3087  
Epoch 20/100  
93/93 — 0s 844us/step - loss: 16.7961 - mae: 17.2883  
Epoch 21/100  
93/93 — 0s 836us/step - loss: 16.7969 - mae: 17.2879  
Epoch 22/100  
93/93 — 0s 833us/step - loss: 16.7633 - mae: 17.2560  
Epoch 23/100  
93/93 — 0s 817us/step - loss: 16.7484 - mae: 17.2407  
Epoch 24/100  
93/93 — 0s 820us/step - loss: 16.7268 - mae: 17.2188  
Epoch 25/100  
93/93 — 0s 840us/step - loss: 16.7147 - mae: 17.2075  
Epoch 26/100  
93/93 — 0s 839us/step - loss: 16.7053 - mae: 17.1969  
Epoch 27/100  
93/93 — 0s 823us/step - loss: 16.6945 - mae: 17.1866  
Epoch 28/100  
93/93 — 0s 824us/step - loss: 16.6705 - mae: 17.1632  
Epoch 29/100  
93/93 — 0s 828us/step - loss: 16.6536 - mae: 17.1456  
Epoch 30/100  
93/93 — 0s 822us/step - loss: 16.6607 - mae: 17.1538  
Epoch 31/100  
93/93 — 0s 823us/step - loss: 16.6098 - mae: 17.1020  
Epoch 32/100  
93/93 — 0s 809us/step - loss: 16.6365 - mae: 17.1284  
Epoch 33/100  
93/93 — 0s 843us/step - loss: 16.6002 - mae: 17.0930  
Epoch 34/100  
93/93 — 0s 829us/step - loss: 16.5956 - mae: 17.0874  
Epoch 35/100  
93/93 — 0s 844us/step - loss: 16.5704 - mae: 17.0621  
Epoch 36/100  
93/93 — 0s 832us/step - loss: 16.5811 - mae: 17.0723  
Epoch 37/100  
93/93 — 0s 828us/step - loss: 16.5902 - mae: 17.0836  
Epoch 38/100  
93/93 — 0s 818us/step - loss: 16.5834 - mae: 17.0777  
Epoch 39/100  
93/93 — 0s 830us/step - loss: 16.5633 - mae: 17.0558  
Epoch 40/100  
93/93 — 0s 840us/step - loss: 16.5631 - mae: 17.0557  
Epoch 41/100  
93/93 — 0s 839us/step - loss: 16.5518 - mae: 17.0449  
Epoch 42/100

93/93 — 0s 819us/step - loss: 16.5367 - mae: 17.0295  
Epoch 43/100  
93/93 — 0s 814us/step - loss: 16.5339 - mae: 17.0261  
Epoch 44/100  
93/93 — 0s 814us/step - loss: 16.5379 - mae: 17.0305  
Epoch 45/100  
93/93 — 0s 831us/step - loss: 16.5371 - mae: 17.0296  
Epoch 46/100  
93/93 — 0s 837us/step - loss: 16.5181 - mae: 17.0100  
Epoch 47/100  
93/93 — 0s 835us/step - loss: 16.5070 - mae: 16.9984  
Epoch 48/100  
93/93 — 0s 850us/step - loss: 16.4951 - mae: 16.9863  
Epoch 49/100  
93/93 — 0s 842us/step - loss: 16.4859 - mae: 16.9767  
Epoch 50/100  
93/93 — 0s 822us/step - loss: 16.4808 - mae: 16.9708  
Epoch 51/100  
93/93 — 0s 830us/step - loss: 16.4734 - mae: 16.9642  
Epoch 52/100  
93/93 — 0s 813us/step - loss: 16.4631 - mae: 16.9539  
Epoch 53/100  
93/93 — 0s 814us/step - loss: 16.4574 - mae: 16.9487  
Epoch 54/100  
93/93 — 0s 788us/step - loss: 16.4784 - mae: 16.9698  
Epoch 55/100  
93/93 — 0s 758us/step - loss: 16.4833 - mae: 16.9756  
Epoch 56/100  
93/93 — 0s 785us/step - loss: 16.4793 - mae: 16.9712  
Epoch 57/100  
93/93 — 0s 779us/step - loss: 16.4487 - mae: 16.9402  
Epoch 58/100  
93/93 — 0s 778us/step - loss: 16.4443 - mae: 16.9351  
Epoch 59/100  
93/93 — 0s 787us/step - loss: 16.4602 - mae: 16.9520  
Epoch 60/100  
93/93 — 0s 780us/step - loss: 16.4429 - mae: 16.9341  
Epoch 61/100  
93/93 — 0s 791us/step - loss: 16.4413 - mae: 16.9319  
Epoch 62/100  
93/93 — 0s 783us/step - loss: 16.4467 - mae: 16.9369  
Epoch 63/100  
93/93 — 0s 805us/step - loss: 16.4324 - mae: 16.9227  
Epoch 64/100  
93/93 — 0s 790us/step - loss: 16.4202 - mae: 16.9110  
Epoch 65/100  
93/93 — 0s 800us/step - loss: 16.4187 - mae: 16.9089  
Epoch 66/100  
93/93 — 0s 789us/step - loss: 16.4603 - mae: 16.9491  
Epoch 67/100  
93/93 — 0s 793us/step - loss: 16.4281 - mae: 16.9166  
Epoch 68/100  
93/93 — 0s 818us/step - loss: 16.4120 - mae: 16.9015  
Epoch 69/100  
93/93 — 0s 822us/step - loss: 16.4089 - mae: 16.8975  
Epoch 70/100  
93/93 — 0s 757us/step - loss: 16.3946 - mae: 16.8837  
Epoch 71/100  
93/93 — 0s 809us/step - loss: 16.3890 - mae: 16.8778  
Epoch 72/100  
93/93 — 0s 819us/step - loss: 16.3791 - mae: 16.8678  
Epoch 73/100  
93/93 — 0s 792us/step - loss: 16.3724 - mae: 16.8615  
Epoch 74/100  
93/93 — 0s 854us/step - loss: 16.3807 - mae: 16.8706  
Epoch 75/100  
93/93 — 0s 861us/step - loss: 16.3975 - mae: 16.8874  
Epoch 76/100  
93/93 — 0s 849us/step - loss: 16.3710 - mae: 16.8597  
Epoch 77/100  
93/93 — 0s 834us/step - loss: 16.3550 - mae: 16.8441  
Epoch 78/100  
93/93 — 0s 843us/step - loss: 16.3564 - mae: 16.8466  
Epoch 79/100  
93/93 — 0s 840us/step - loss: 16.3787 - mae: 16.8690  
Epoch 80/100  
93/93 — 0s 825us/step - loss: 16.3696 - mae: 16.8583  
Epoch 81/100  
93/93 — 0s 793us/step - loss: 16.3645 - mae: 16.8528  
Epoch 82/100  
93/93 — 0s 784us/step - loss: 16.3379 - mae: 16.8262  
Epoch 83/100  
93/93 — 0s 783us/step - loss: 16.3288 - mae: 16.8188

```

Epoch 84/100
93/93 ————— 0s 802us/step - loss: 16.3648 - mae: 16.8550
Epoch 85/100
93/93 ————— 0s 809us/step - loss: 16.3548 - mae: 16.8430
Epoch 86/100
93/93 ————— 0s 789us/step - loss: 16.3176 - mae: 16.8072
Epoch 87/100
93/93 ————— 0s 831us/step - loss: 16.3518 - mae: 16.8410
Epoch 88/100
93/93 ————— 0s 838us/step - loss: 16.3571 - mae: 16.8461
Epoch 89/100
93/93 ————— 0s 839us/step - loss: 16.3347 - mae: 16.8224
Epoch 90/100
93/93 ————— 0s 829us/step - loss: 16.3276 - mae: 16.8152
Epoch 91/100
93/93 ————— 0s 823us/step - loss: 16.3045 - mae: 16.7935
Epoch 92/100
93/93 ————— 0s 831us/step - loss: 16.3135 - mae: 16.8015
Epoch 93/100
93/93 ————— 0s 855us/step - loss: 16.3038 - mae: 16.7937
Epoch 94/100
93/93 ————— 0s 826us/step - loss: 16.2962 - mae: 16.7857
Epoch 95/100
93/93 ————— 0s 795us/step - loss: 16.2995 - mae: 16.7895
Epoch 96/100
93/93 ————— 0s 790us/step - loss: 16.2959 - mae: 16.7840
Epoch 97/100
93/93 ————— 0s 801us/step - loss: 16.2939 - mae: 16.7813
Epoch 98/100
93/93 ————— 0s 803us/step - loss: 16.3019 - mae: 16.7888
Epoch 99/100
93/93 ————— 0s 800us/step - loss: 16.3015 - mae: 16.7902
Epoch 100/100
93/93 ————— 0s 797us/step - loss: 16.3021 - mae: 16.7899

```

## Model Prediction

Now see if the model generates good results. If you used the default parameters of this notebook, you should see the predictions follow the shape of the ground truth with an MAE of around 15.

```

In [14]: def model_forecast(model, series, window_size, batch_size):
          """Uses an input model to generate predictions on data windows

          Args:
            model (TF Keras Model) - model that accepts data windows
            series (array of float) - contains the values of the time series
            window_size (int) - the number of time steps to include in the window
            batch_size (int) - the batch size

          Returns:
            forecast (numpy array) - array containing predictions
          """

          # Generate a TF Dataset from the series values
          dataset = tf.data.Dataset.from_tensor_slices(series)

          # Window the data but only take those with the specified size
          dataset = dataset.window(window_size, shift=1, drop_remainder=True)

          # Flatten the windows by putting its elements in a single batch
          dataset = dataset.flat_map(lambda w: w.batch(window_size))

          # Create batches of windows
          dataset = dataset.batch(batch_size).prefetch(1)

          # Get predictions on the entire dataset
          forecast = model.predict(dataset, verbose=0)

          return forecast

```

```

In [15]: # Reduce the original series
          forecast_series = series[split_time-window_size:-1]

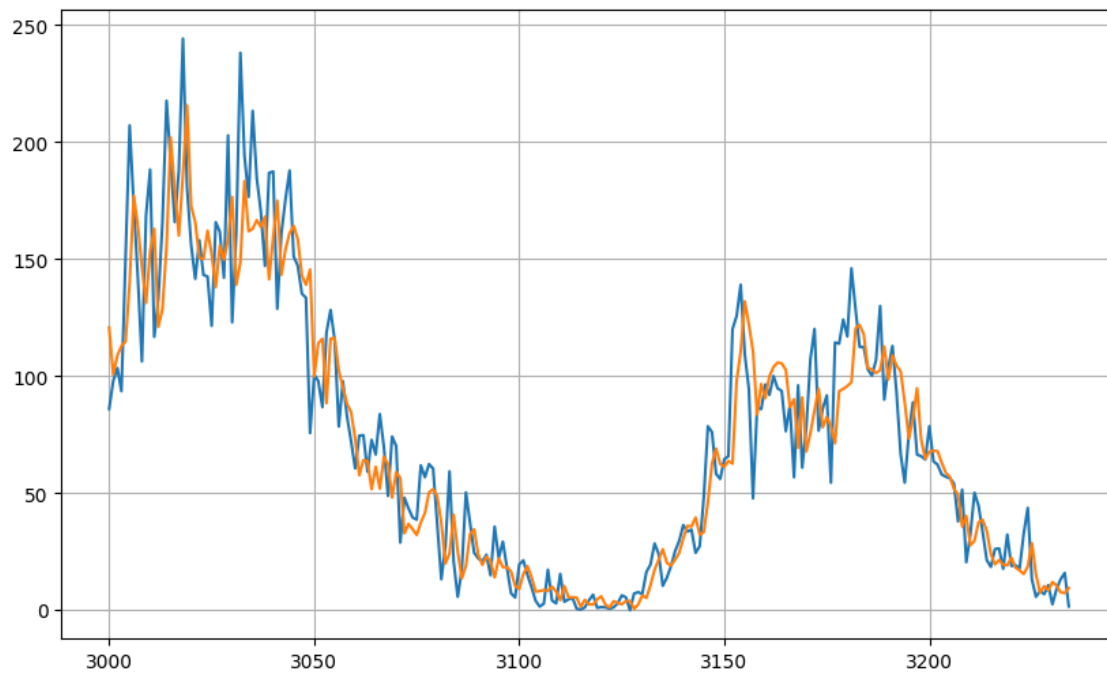
          # Use helper function to generate predictions
          forecast = model_forecast(model, forecast_series, window_size, batch_size)

          # Drop single dimensional axis
          results = forecast.squeeze()

          # Plot the results

```

```
plot_series(time_valid, (x_valid, results))
```



```
In [16]: # Compute the MAE
print(tf.keras.metrics.mae(x_valid, results).numpy())
```

14.920012

## Wrap Up

In this lab, you built a relatively simple DNN to forecast sunspot numbers for a given month. We encourage you to tweak the parameters or train longer and see the best results you can get. In the next lab, you will build a more complex model and you evaluate if the added complexity translates to better or worse results.